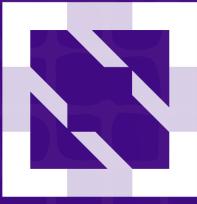




KubeCon



CloudNativeCon

OPEN SOURCE SUMMIT

China 2019





KubeCon



CloudNativeCon



OPEN SOURCE SUMMIT

China 2019

7 Missing Factors for Your Production-Quality 12-Factor Apps

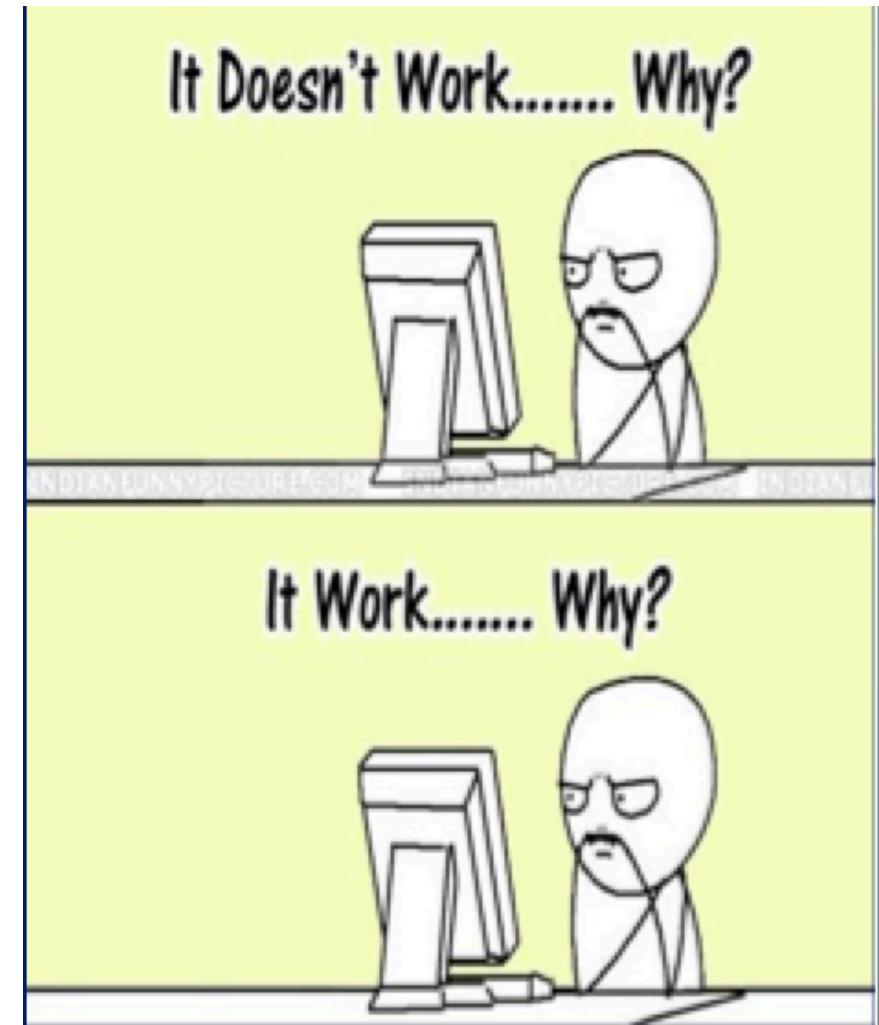
Michael Elder
IBM Distinguished Engineer - IBM Multicloud Platform
@mdelder

Shikha Srivastava
IBM Senior Technical Staff Member
@shikhasthoughts



What is Ready for production application

- Secure
 - Installation, authentication and access
- Resilient, Highly Available and scale
- Repeated deployment
 - with safe upgrades and configuration changes
- Performance
- Observable
- Upgradeable
- more
- And AGILE too



- “12-Factor” is a software methodology for building scalable microservice applications
- Originally created by Heroku
- Best practices designed to enable applications to be built with ***portability, resilience, and scalability*** when deployed to the web

What is a 12-factor app?



kubernetes & 12-factor apps

Why

12 factor apps?

I. Codebase

One codebase tracked in revision control, many deploys

II. Dependencies

Explicitly declare and isolate dependencies

III. Config

Store config in the environment

IV. Backing services

Treat backing services as attached resources

V. Build, release, run

Strictly separate build and run stages

VI. Processes

Execute the app as one or more stateless processes

VII. Port binding

Export services via port binding

VIII. Concurrency

Scale out via the process model

IX. Disposability

Maximize robustness with fast startup and graceful shutdown

X. Dev/prod parity

Keep development, staging, and production as similar as possible

XI. Logs

Treat logs as event streams

XII. Admin processes

Run admin/management tasks as one-off processes

- Make it easier to run, scale, and deploy applications
- Keep parity between development and production
- Provide strict separation between build, release, and run stages

Code

I. Codebase

One codebase tracked in revision control, many deploys

V. Build, release, run

Strictly separate build and run stages

X. Parity between dev & prod

Keep development, staging, and production as similar as possible

Deploy

II. Dependencies

Explicitly declare and isolate dependencies

III. Config

Store config in the environment

IV. Backing services

Treat backing services as attached resources

VI. Processes

Execute the app as one or more stateless processes

VII. Port binding

Export services via port binding

Operate

VIII. Concurrency

Scale out via the process model

IX. Disposability

Maximize robustness with fast startup and graceful shutdown

XI. Logs

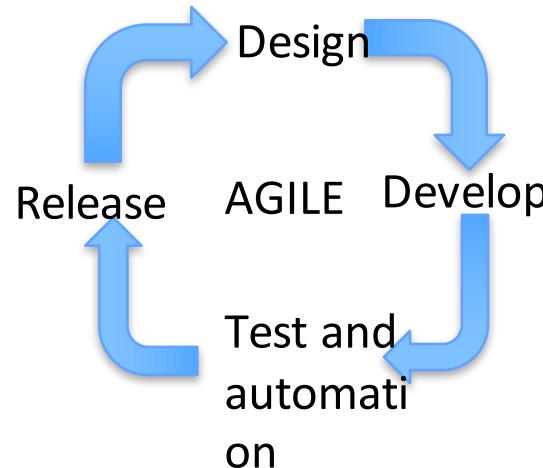
Treat logs as event streams

XII. Admin processes

Run admin/management tasks as one-off processes

Developers dream – Code factors

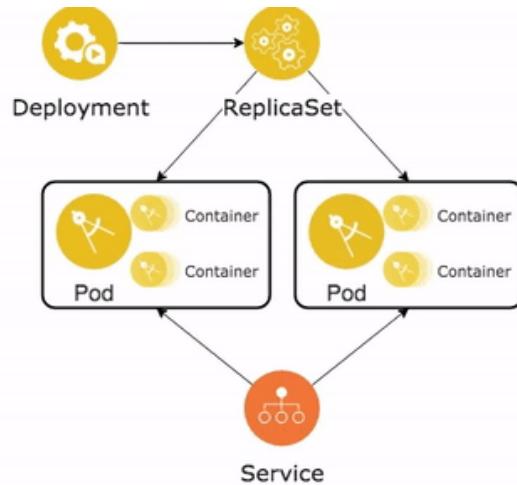
- One codebase for my application tracked in revision that runs anywhere: build, ship and run anywhere
- AND
- I can offload deployment, HA, scaling, upgrade strategy and not worry about it



- **Container Images** built from Dockerfiles using trusted small image. **Kubernetes Deployments, etc** managed as YAML (F#I- *Codebase*)
- Having a strong artifact-driven model makes it easier to follow a **Continuous Delivery** lifecycle (F#V- *Build, release, run*)
- Using the same **images** and YAML objects make it easier for **dev teams** to match what's running in **production** (F#X- *Dev/prod parity*)

Operate factors: Concurrency (F#VIII) & Disposability (F#IX)

- Ensure scale for your app
- Replica set ensures specified number of pods are always running



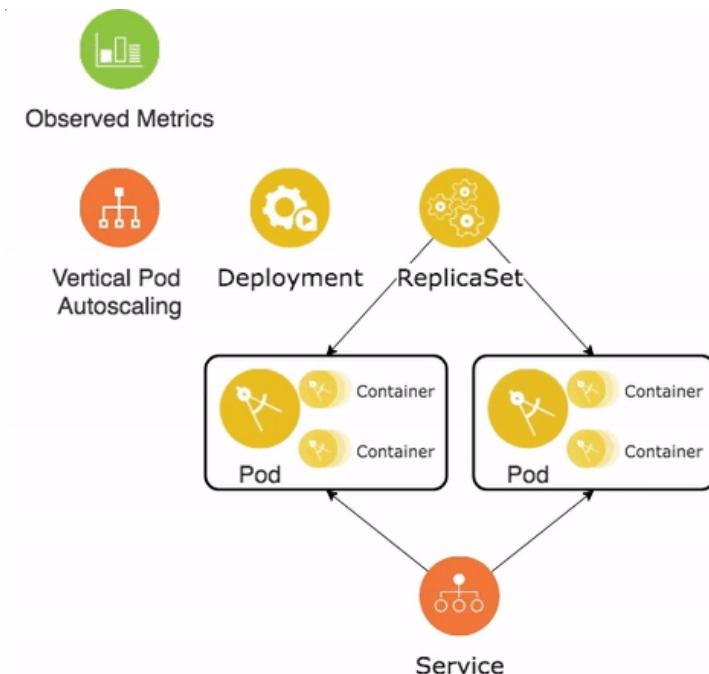
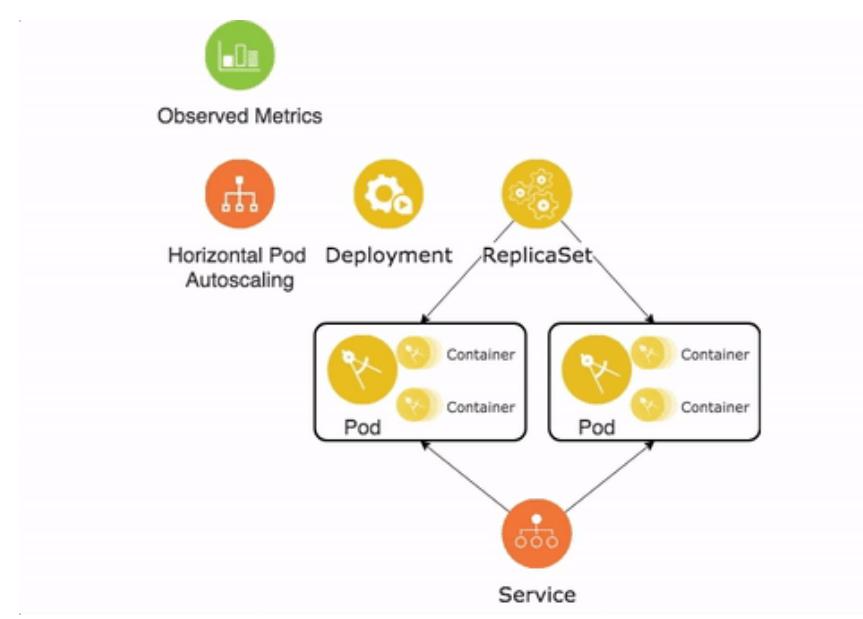
- Is this enough?
Remember load is never constant in the real world

```
kind: Deployment
metadata:
  name: nginx
spec:
  replicas: 2
  template:
    metadata:
      labels:
        service: http-server
    spec:
      containers:
      - name: nginx
        image: nginx:1.10.2
        imagePullPolicy: IfNotPresent
      ports:
      - containerPort: 80
```

Operate factors: Concurrency (F#VIII)

Leverage autoscaling to automate computation resources based on load

- Horizontal Pod Scaler (HPA)
 - Controls the number of replicas
 - Use cpu or memory as a trigger or use custom metric
 - Applicable for stateless app
- Vertical Pod Scaler (VPA)
 - Controls the memory and cpu for pod
 - Use cpu or memory as a trigger or use custom metric
 - Applicable for statefull apps



7

missing factors

XIII. Observable

Apps should provide visibility about current health and metrics

XIV. Schedulable

Apps should provide guidance on expected resource constraints

XV. Upgradable

Apps must upgrade data formats from prior generations

XVI. Least privileged

Apps should provide guidance on expected resource constraints

XVII. Auditable

Apps should provide appropriate audit logs for compliance needs

XVIII. Access Control (Identity, Network, Scope, Certificates)

Protect app and resources from the world

XIX. Measurable

Apps usage should be measurable for quota or chargebacks



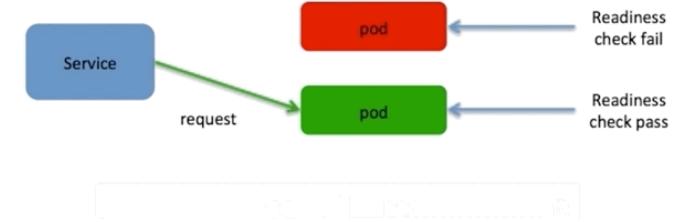
7 missing factors from 12 factor application

Observable: Application health (F#XIII)

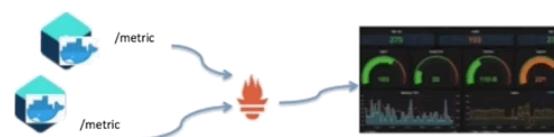
Know your application health

- Kubernetes probes
 - Is the app ready to accept traffic?: Readiness
 - Is the app responsive? : Liveliness
- Is this enough?
 - What about transactions, traffic, memory usage ?

```
readinessProbe:  
  # an http probe  
  httpGet:  
    path: /readiness  
    port: 8080  
    initialDelaySeconds:  
      20  
    periodSeconds: 5
```



```
livenessProbe:  
  # an http probe  
  httpGet:  
    path: /healthcheck  
    port: 8080  
    initialDelaySeconds: 15  
    timeoutSeconds: 1
```





KubeCon



CloudNativeCon

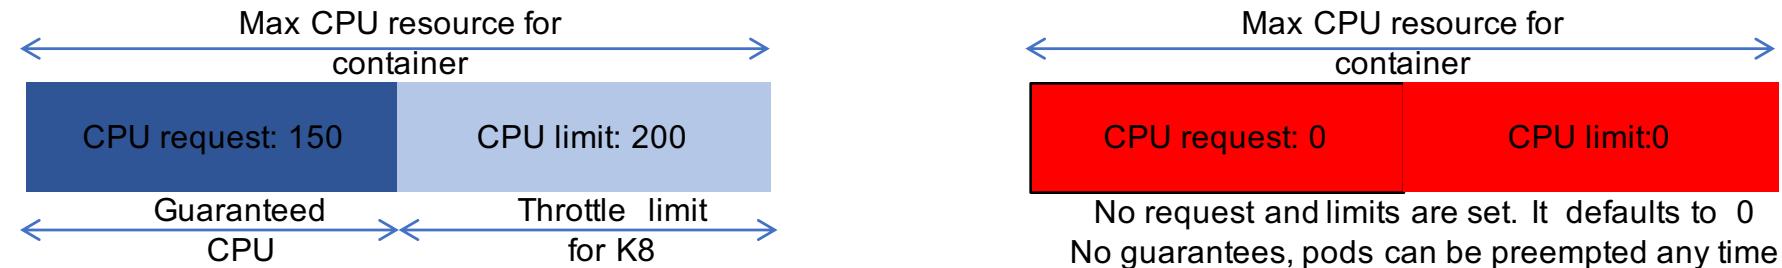


OPEN SOURCE SUMMIT

China 2019

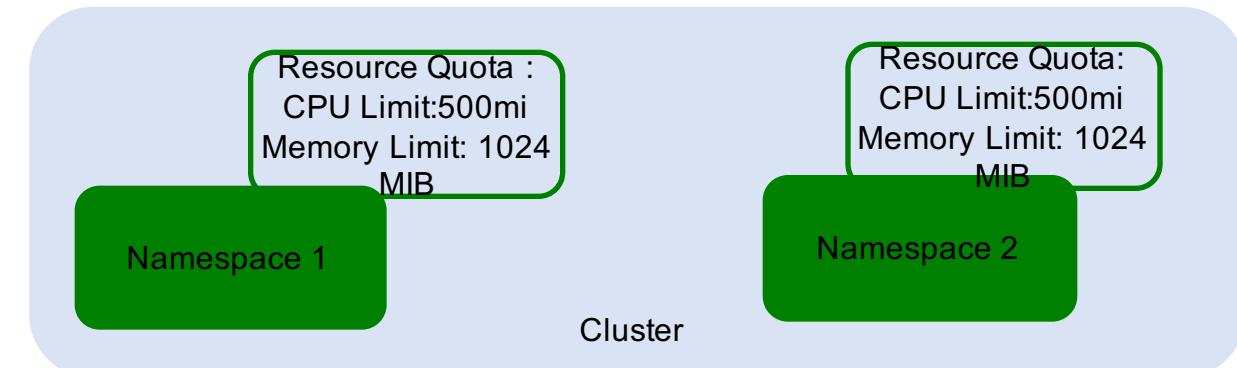
Schedulable: Resource requests, limits, & quotas (F#XIV)

Guarantee resources for your containers: Specify request and limits for the compute resources



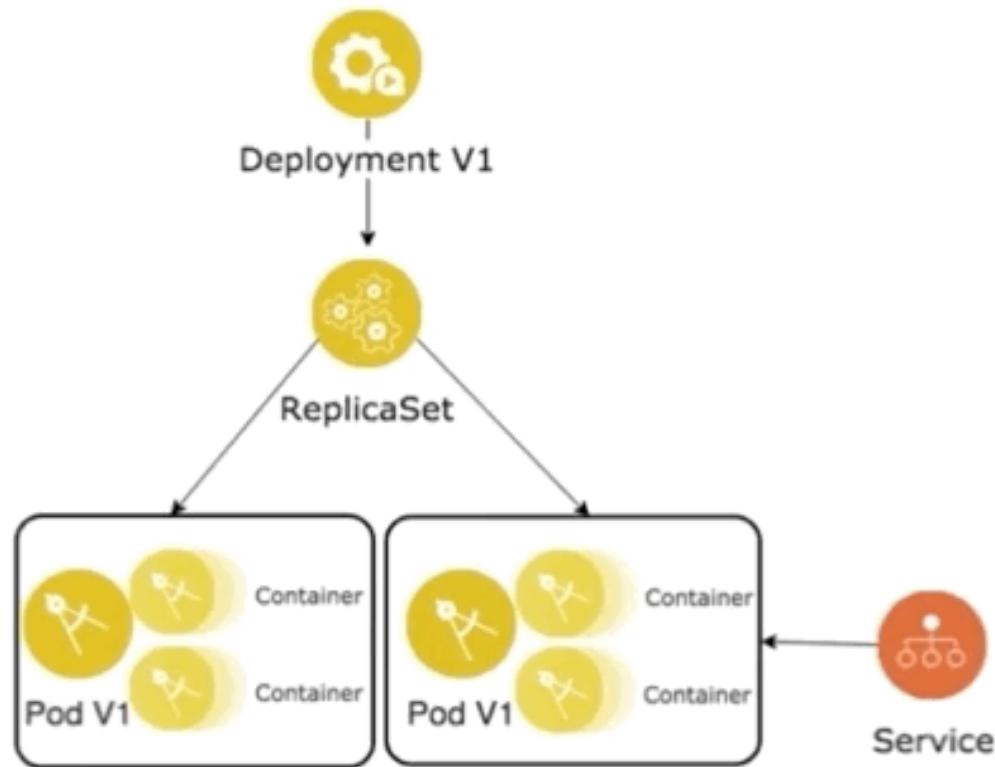
Set resource quota

Once quota in a namespace for compute resources set, the users are forced to set requests or limits for those values



Upgradable (F#XV)

Applications should be able to roll out updates for cases where backward compatible updates (security or feature updates)needs to be made

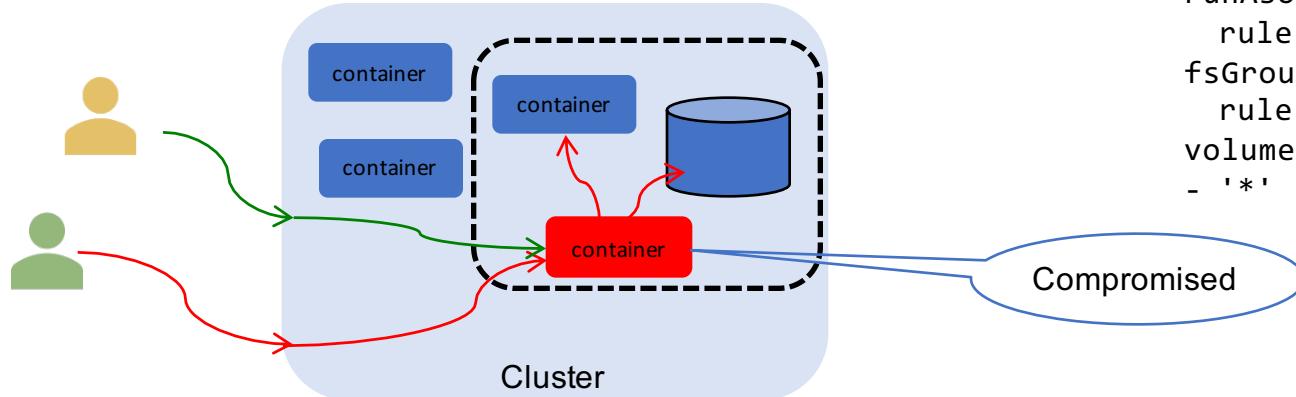


```
minReadySeconds: 5
strategy:
  # indicate which
  strategy
  # we want for rolling
  update
  type: RollingUpdate
  rollingUpdate:
    maxSurge: 1
    maxUnavailable: 1
```

Least Privilege (F#XVI)

Limit container access to hosts. Every permission is an attack vector

- Use Pod Security Policy and Network Policy to
 - Limit access to filesystem
 - Limit access to Kernel capabilities
 - Use a non-privileged user
 - Limit access to volume types
 - Limit access to ports



```
#sample-psp.yaml
apiVersion: policy/v1beta1
kind: PodSecurityPolicy
metadata:
  name: example
spec:
  privileged: false
  # Don't allow
  #   privileged pods!
  # The rest fills in some
  #   required fields.
  seLinux:
    rule: RunAsAny
  supplementalGroups:
    rule: RunAsAny
  runAsUser:
    rule: RunAsAny
  fsGroup:
    rule: RunAsAny
  volumes:
  - '*'
```

Auditable (F#XVII)

- Know WHAT/WHEN/WHO/WHERE for all CRUD operations
 - Chronological set of records documenting sequence of events affecting system and application by users or components
- Use cloud agnostic industry standard format – CADF (Cloud Auditing Data Federation)
- Control the quantity of logs

CADF event:

<initiator_id>: ID of the user that performed the operation

<target_uri>: CADF specific target URI, (for example: data/security/project)

<action>: The action being performed, typically: <operation>. <resource_type>

Access Control -Identity, Network, Scope (F#XVIII)

Protect app and resources from the world

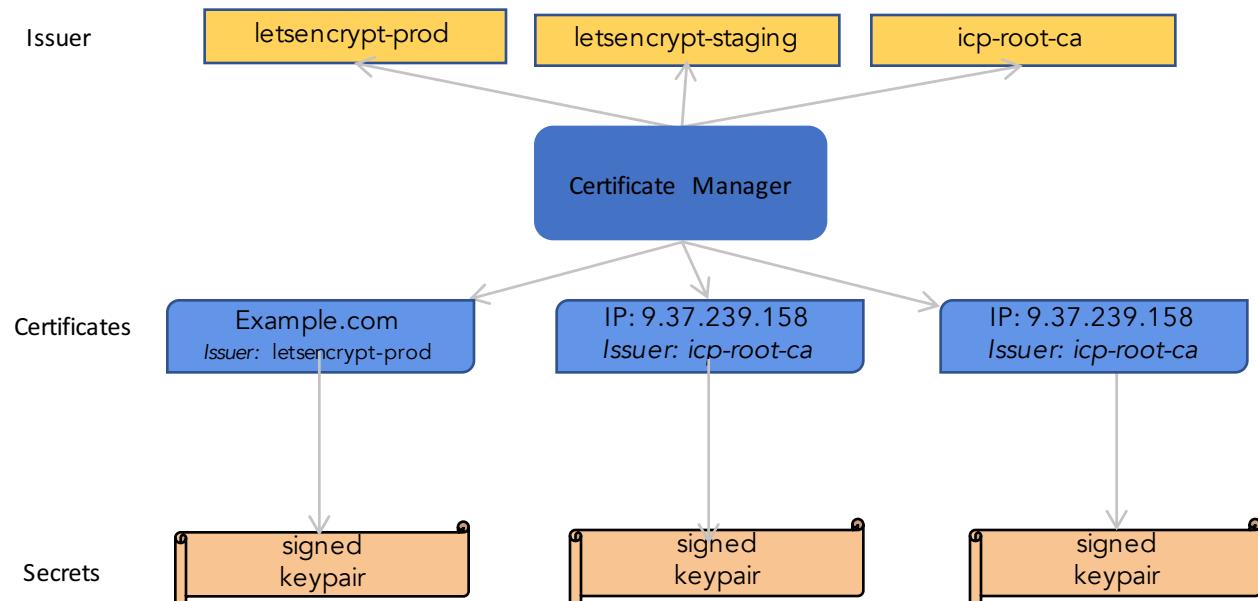
- Authentication and Authorization
- Certificate Management
- Data Protection
- Network security
 - Network policy
 - Network Isolation
- Admission Controller
 - Example: Image admission controller



Access Control: Identity, Network, Scope (F#XVIII)

Ensure secure communication

- Generate Certificates
- Enable TLS / mTLS
- Manage Certificates



```
# sample issuer.yaml
apiVersion: certmanager.k8s.io/v1alpha1
kind: Issuer
metadata:
  name: demo1-nginx-ca
  namespace: demo
spec:
  ca:
    secretName: demo1-nginx-ca-key-pair
```

1. Issuer creates Certificate

```
# sample certificate.yaml
apiVersion: certmanager.k8s.io/v1alpha1
kind: Certificate
Metadata:
  name: demo1-nginx-cert
spec:
  secretName: demo1-nginx-cert
  issuerRef:
    name: demo1-nginx-ca
    kind: Issuer
    commonName: "foo1.bar"
    dnsNames:
      - foo1.bar1
```

2. Certificate creates secret

```
ls -l /var/lib/etcd/member/secrets/namespaces/demo/secrets/demo1-nginx-ca-key-pair
Name: demo1-nginx-ca-key-pair
Namespace: demo
Labels: <none>
Annotations: <none>

Type: kubernetes.io/tls
Data
=====
tls.crt: 1598 bytes
tls.key: 1679 bytes
admin:-$
```

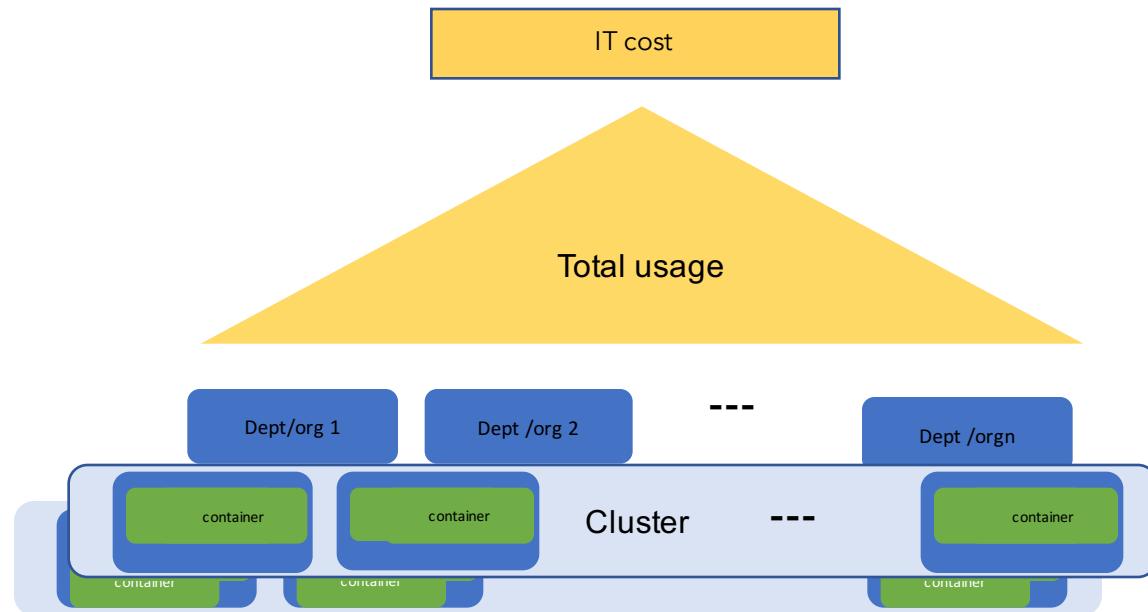
3. Secret mounts to Pod

Pod

Measurable (F#XIX)

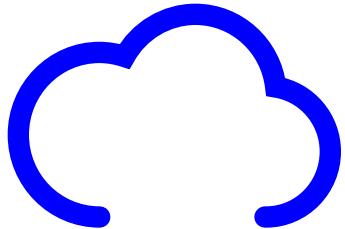
Know the cost of the application

- Compute resources allocated to run the containers should be measurable
- Org / department using the cluster should be accountable



So, what really
makes a production-
ready app?

A production grade application



Attention to
Cloud provider configurations

Example XII: Observable, Example: XVIII: Access Control.
Factor XIX:: Measurable

Attention to
Kubernetes configuration

Example: Factor III: Config, Factor II Config, Factor XIV:
Schedulable

Attention to
Building containers and what's inside the containers

Example: Factor I : codebase , Factor X: dev/prod parity,
Factor XV

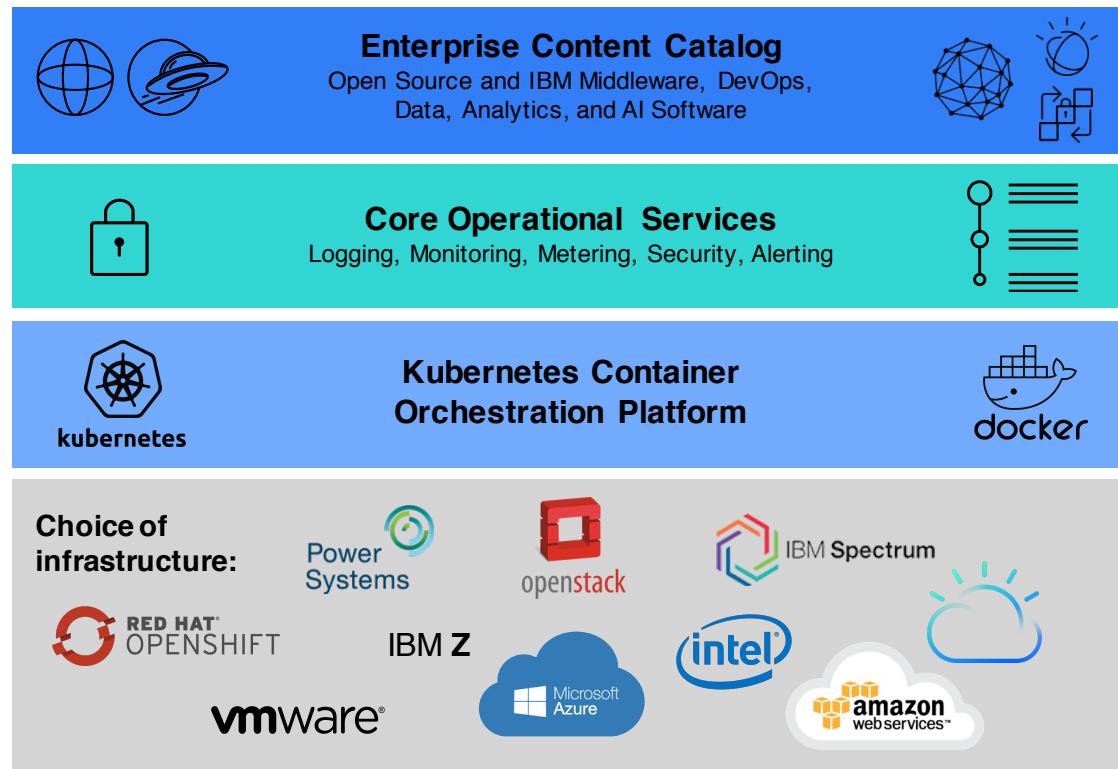


Production thinking needs to be through the entire process



Enough
talking, let's
see it LIVE!

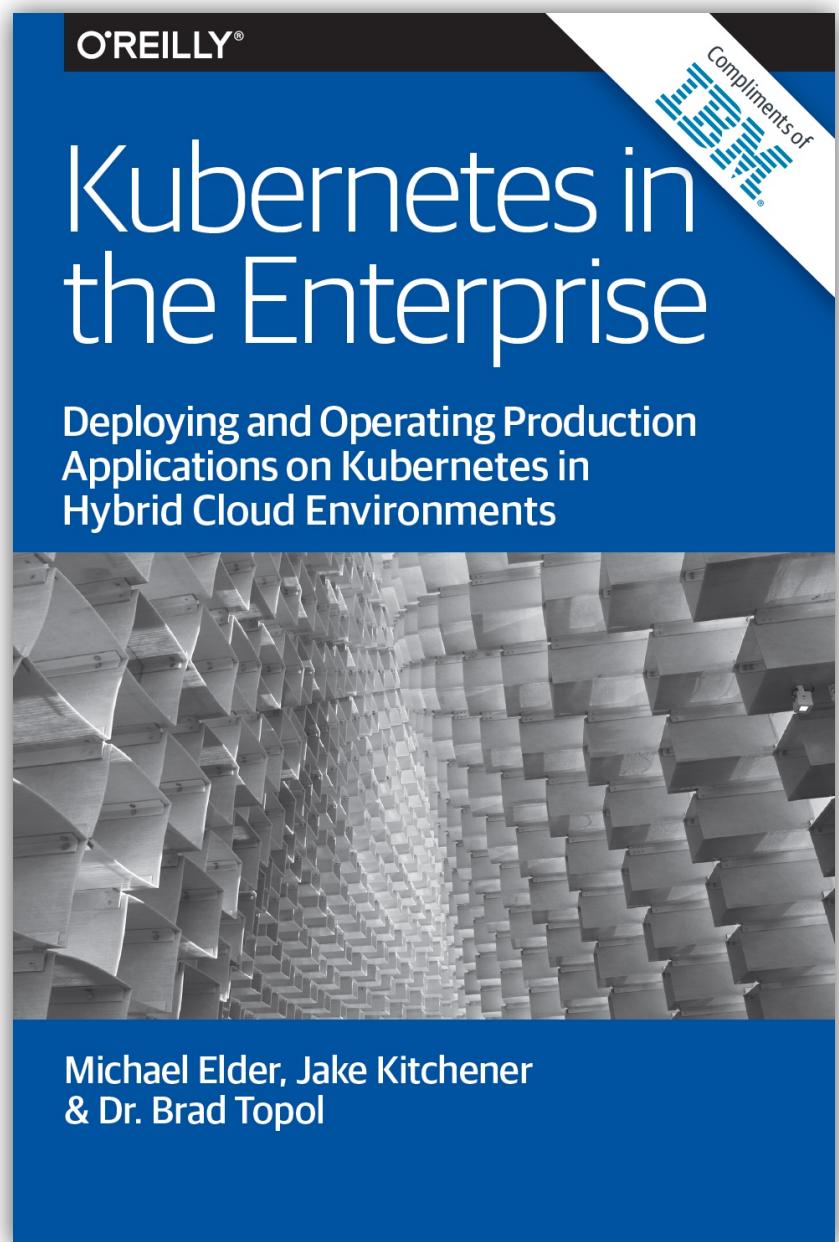
IBM Cloud Private (ICP)



Provides the capabilities to run containerized application in secure, scalable and resilient environment

- Self-service rich catalog of IBM MW
- Helm based parameterized install to simplify complex K8 apps
- Logging : ELK + filebeat
- Monitoring : Prometheus + Grafana
- Usage : IBM Metering Service
- IBM Vulnerability Advisor
- IBM Mutation Advisor
- Authentication/ Authorization
- Certificate Management
- Network security
- Audit trail for any CRUD operations
- Team based organization of resources

All communication enabled over TLS. Data secured in transit and at rest



Now available online compliments of IBM: ibm.biz/BdYA4i >

Need Details on signing

#7678A: Tech Talk:
Deploying Kubernetes in
the Enterprise (with the
authors)

When: Wednesday, 11:30
AM - 12:10 PM

Where: Table Top Tap Room
at the Metreon | Code Cafe
Tech Talks Area

Get a **signed** copy with all
of the authors at the Code
Café Mezzanine on
Wednesday (7 – 7:30PM)!

Learn
more in
our new
book!