



Network Observability with Envoy

Network Observability with Envoy

Nic Jackson

Developer Advocate at HashiCorp





Agenda

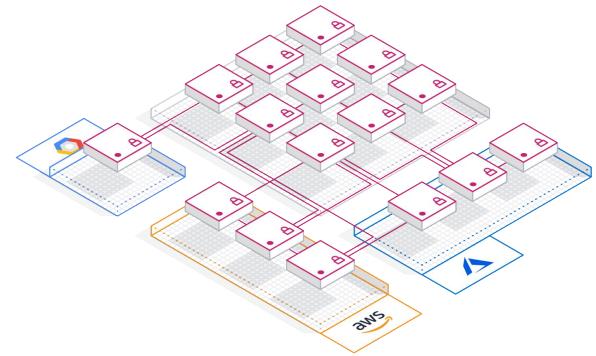
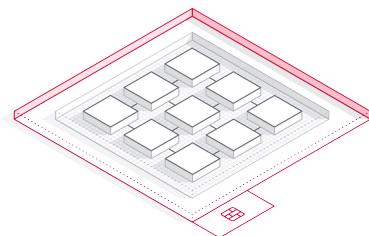
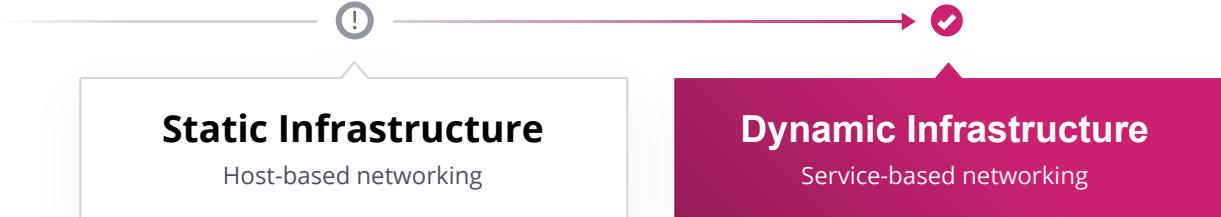
1. **Introduction** - create a common vocabulary.
2. **Metrics** - configure and use metrics.
3. **Tracing** - configure and use tracing.
4. **Logging** - collect access logs.



Introduction



The shift from static to dynamic networking





The shift from static to dynamic networking



Static Infrastructure

Host-based networking

Homogeneous infrastructure with static IPs, primarily north-south traffic, protected by coarse-grained segments.

TRADITIONAL APPROACH

- Load balancers to create artificial static IPs
- Firewall sprawl to constrict service traffic
- Configuration management to deploy services

Dynamic Infrastructure

Service-based networking

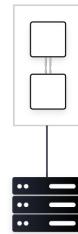
Heterogeneous infrastructure with dynamic IPs, dominated by east-west traffic, without a clear network perimeter.

CONSUL APPROACH

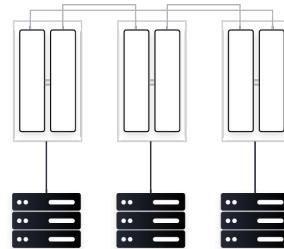
- Service Discovery for connectivity
- Service Segmentation for security
- Service Configuration for runtime configuration



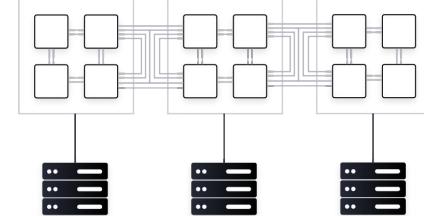
Market trend from **monoliths** to **microservices**



Single,
Physical
Server



Dynamic Virtual
Machines



Smaller,
Ephemeral
Containers



Business challenges of dynamic infrastructure



Reduced Productivity

Waiting for manual updates to load balancers and firewalls blocks development throughput.



Increased Risk

Firewall rule sprawl is complex to manage and mistakes create security vulnerabilities.



Increased Cost

Load balancers and firewalls are expensive and costly to maintain.

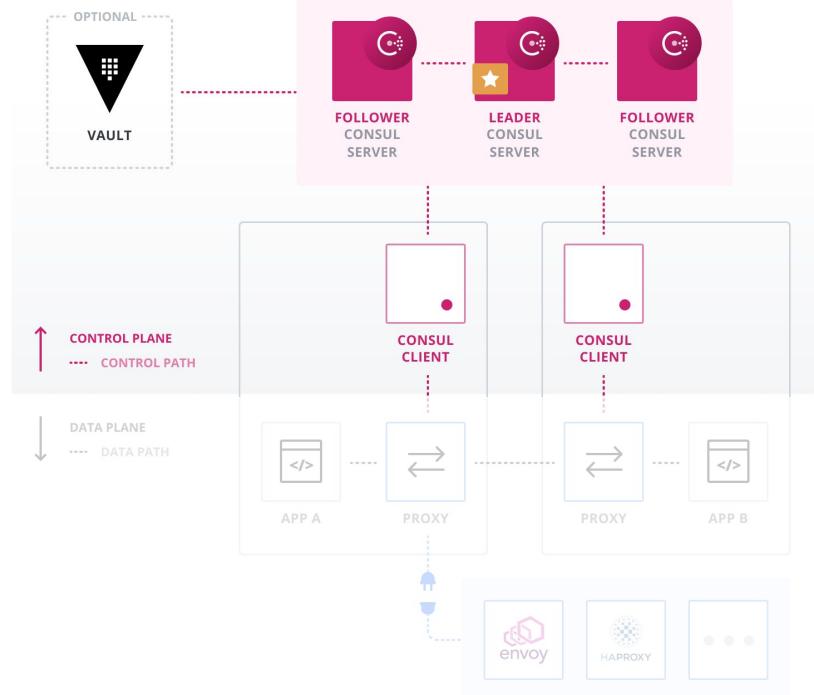
Components



Service Mesh

Control Plane

- Service to service communication policy
- Service Catalog
- CA and x509 certificate generation
- Configuration and proxy management

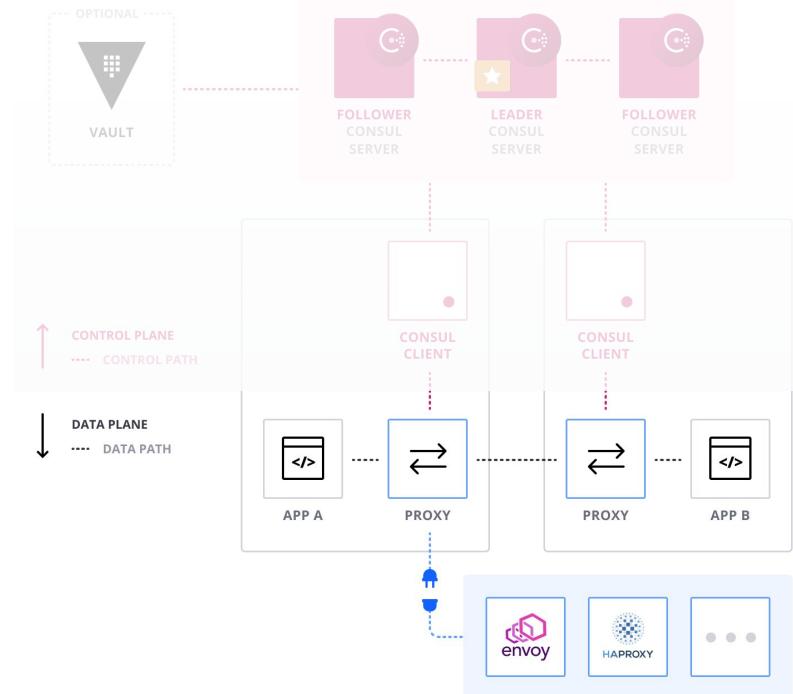


Components

Service Mesh

Data Plane

- Authorization
- Request tracing
- Traffic shaping
- Load balancing
- Service discovery
- Circuit breaking
- Retry logic
- Networking statistics



Networks are **not** 100% stable and
often experience transient failure.



You can't do **Reliability** without
Observability.



Observability, is it **just a buzzword?**

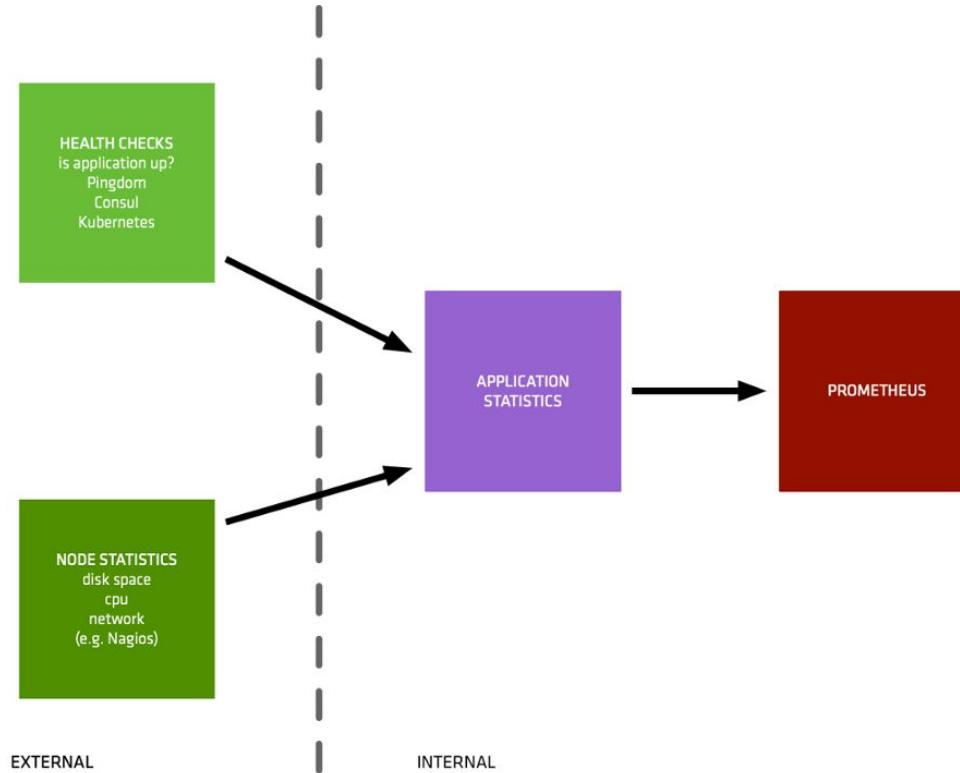
Observability is a measure of how well
internal states of a system can be
inferred from knowledge of
external outputs.



Observability



Internal and external instrumentation





—

Metrics



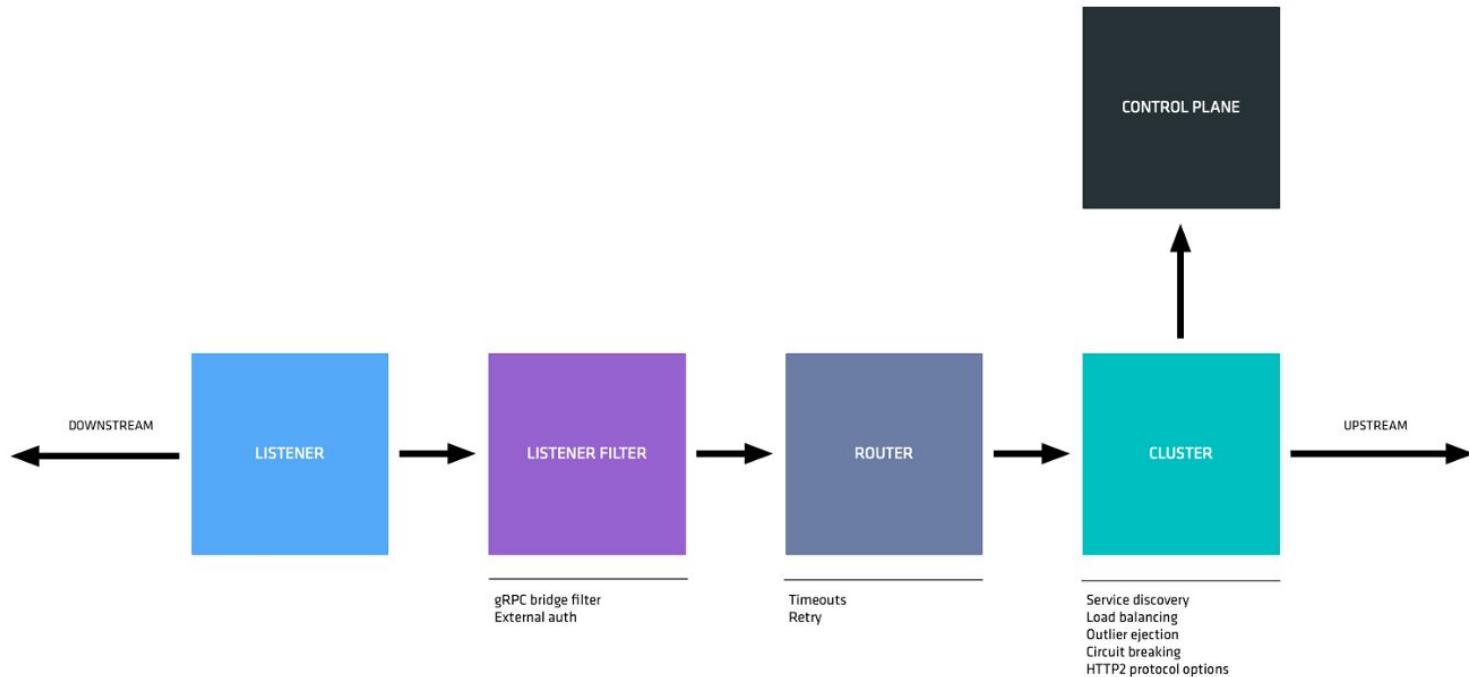
—

Envoy Architecture

Metrics



Terminology

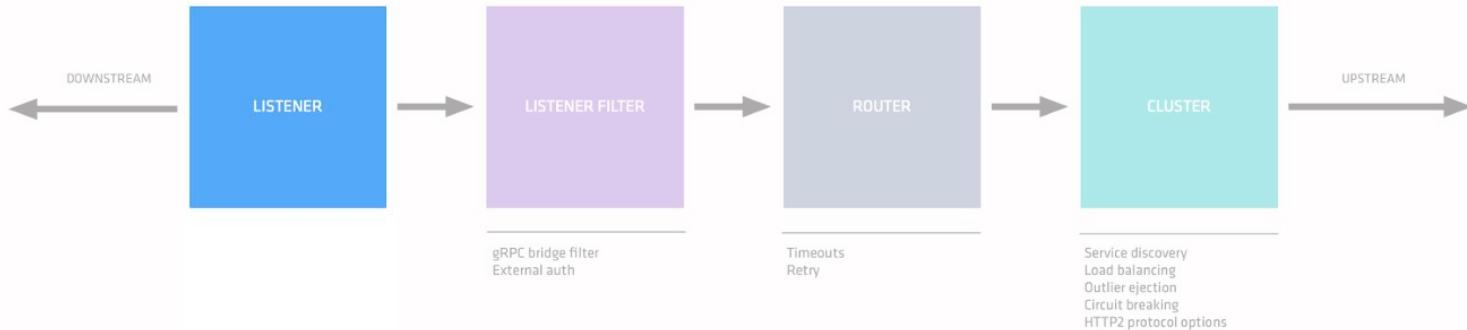




Listener

Terminology

A listener is a named network location (e.g., port, unix domain socket, etc.) that can be connected to by downstream clients. Envoy exposes one or more listeners that downstream hosts connect to.

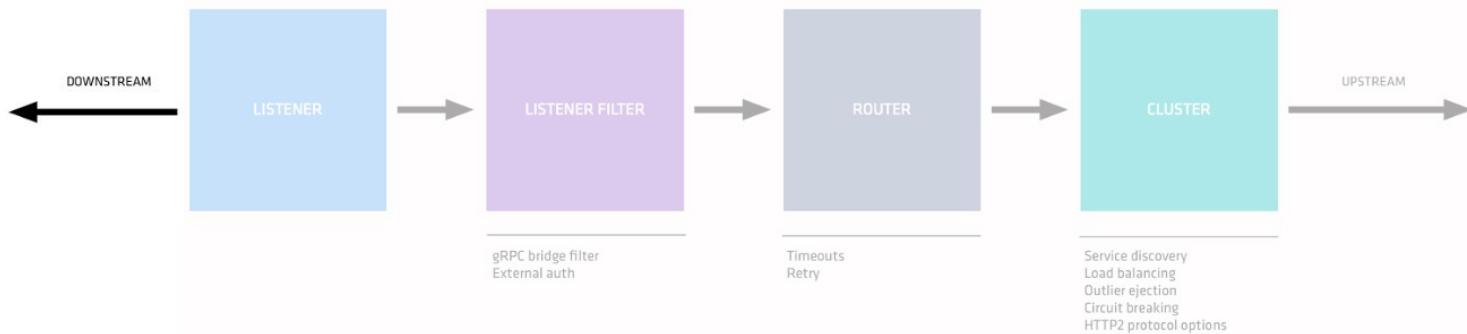




Downstream

Terminology

A downstream host connects to Envoy, sends requests, and receives responses.

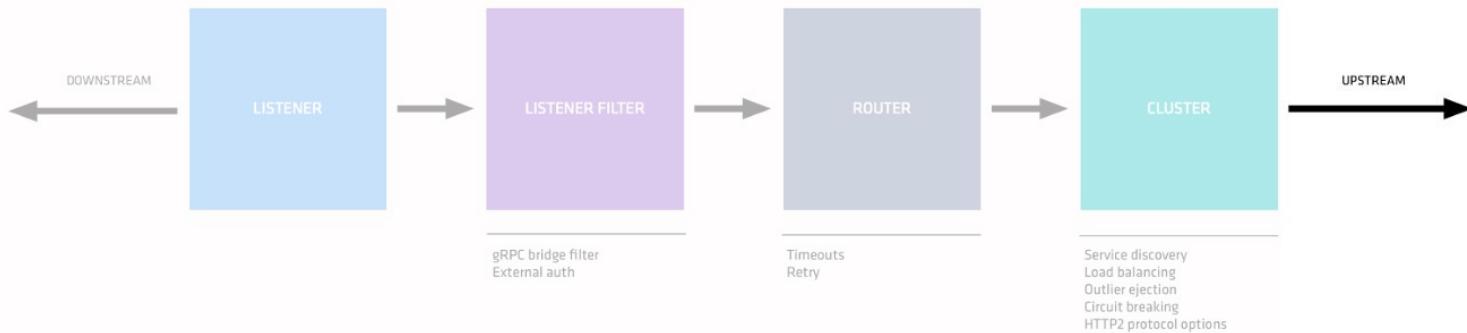




Upstream

Terminology

An upstream host sends requests from Envoy to other services and returns responses.

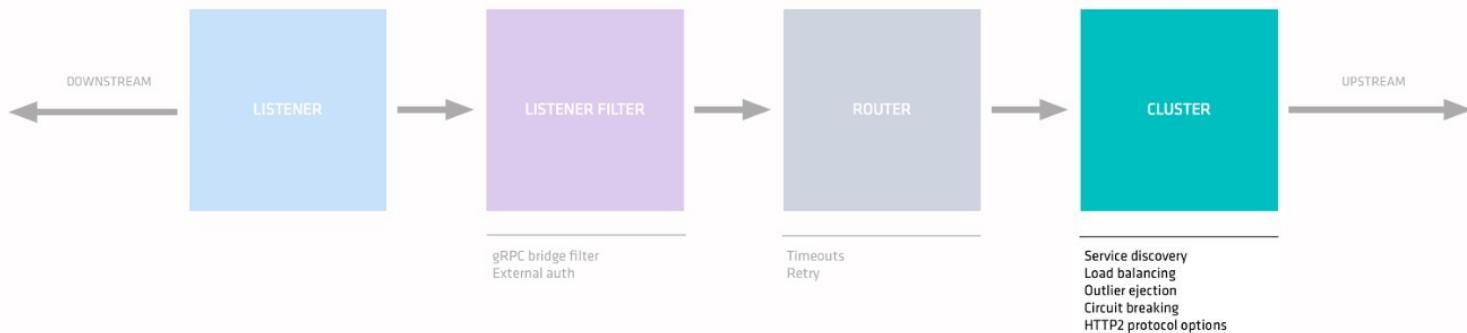




Cluster

Terminology

A cluster is a group of logically similar upstream hosts that Envoy connects to. Envoy discovers the members of a cluster via **service discovery**. The cluster member that Envoy routes a request to is determined by the **load balancing policy**.





—

Configuration

Metrics



Envoy metrics.

Envoy does not label the metrics with the application name, so **add tags** to be able **to differentiate** between metrics.

CODE EDITOR

```
"stats_config": {  
    "stats_tags": [  
        {  
            "tag_name": "local_cluster",  
            "fixed_value": "emojify-api-v2"  
        }  
    ],  
    "use_all_default_tags": true  
}
```



Envoy Prometheus Metrics

Metrics

- **1.10** introduces histograms for Prometheus metrics
- Metrics exposed with **unsecured** admin endpoint (`/stats/prometheus`),
- Exposure of metrics **needs** to be configured with **loopback** route to
avoid exposing **admin endpoints**



—

Using metrics

Metrics



StatsD

- Originally created by Etsy
- Push based metrics
- Lightweight UDP protocol
- **No support for metadata**



Metrics types

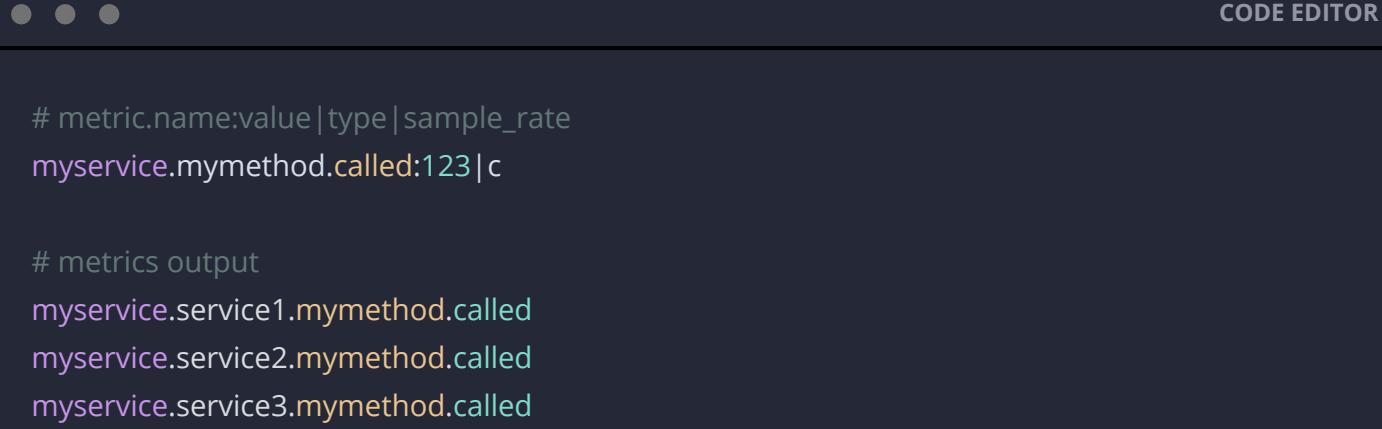
StatsD

Type	Description
Counter	Increment value, e.g. number of method calls.
Gauge	Value over time, e.g. CPU consumption, memory usage.
Timing	Time taken to perform a task, e.g. time take to perform a method call.
Set	Set of unique values over collection period.



Metrics format

StatsD does not support basic metric labels.



The image shows a dark-themed code editor window with three dots at the top left and the text "CODE EDITOR" at the top right. The editor contains the following text:

```
# metric.name:value|type|sample_rate
myservice.mymethod.called:123|c

# metrics output
myservice.service1.mymethod.called
myservice.service2.mymethod.called
myservice.service3.mymethod.called
```



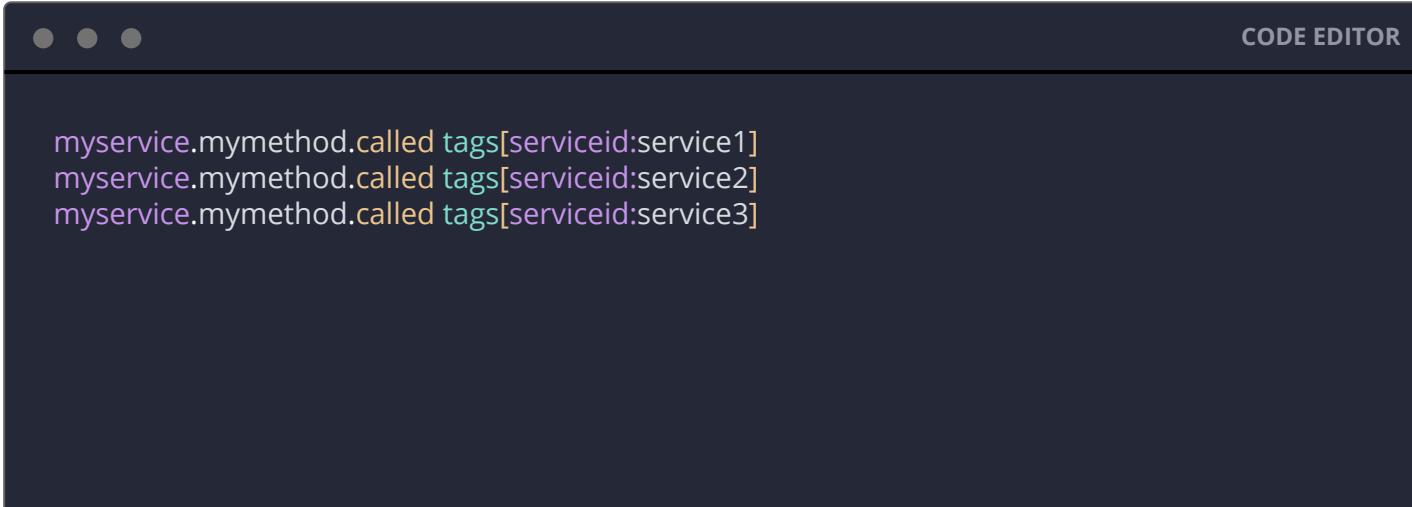
DogStatsD

- Created by DataDog based on StatsD protocol
- Push based metrics
- Lightweight UDP protocol
- **Support for metadata through tags**



Metrics format

DogStatsD



The image shows a dark-themed code editor window. At the top, there are three small circular navigation icons on the left and the text "CODE EDITOR" on the right. The main area contains three lines of text, each consisting of a metric name followed by ".called" and then "tags[serviceid:service1]" in purple. The lines are:

```
myservice.mymethod.called tags[serviceid:service1]
myservice.mymethod.called tags[serviceid:service2]
myservice.mymethod.called tags[serviceid:service3]
```



Prometheus

- Pull based approach from central server
- Service implements HTTP endpoint exposing metrics
- **Supports metadata by default**



Metrics types

Prometheus

Type	Description
Counter	Cumulative metric, representing a monotonically increasing counter, e.g. number of method calls.
Gauge	Single numerical value that can arbitrarily go up and down, e.g. CPU consumption.
Histogram	Samples observations and counts them in configurable buckets, e.g. request timings.



Metrics format

Prometheus

A screenshot of a dark-themed code editor window. At the top, there are three small circular icons on the left and the text "CODE EDITOR" on the right. The main area contains a single line of text: "envoy_http_downstream_rq_completed{envoy_http_conn_manager_prefix="ingress_cache"}".

```
envoy_http_downstream_rq_completed{envoy_http_conn_manager_prefix="ingress_cache"}
```



Choosing a format

- Tagging is **essential** to effectively build dashboards
- Metrics **need** to be tagged with **Metadata** such as pod name, node, etc



—

Listener

Metrics



Key service metrics

Listener - Connections

Every listener has a statistics tree rooted at <prefix>.listener.<address>. with the following statistics:

downstream_cx_total	Counter	Total connections
downstream_cx_destroy	Counter	Total destroyed connections
downstream_cx_active	Gauge	Total active connections



Envoy metrics.

use_all_default_tags

extracts **common**
components from metric
names and **adds as tags**

CODE EDITOR

```
"stats_config": {  
    "stats_tags": [  
        {  
            "tag_name": "local_cluster",  
            "fixed_value": "emojify-api-v2"  
        }  
    ],  
    "use_all_default_tags": true  
}
```



Metrics queries

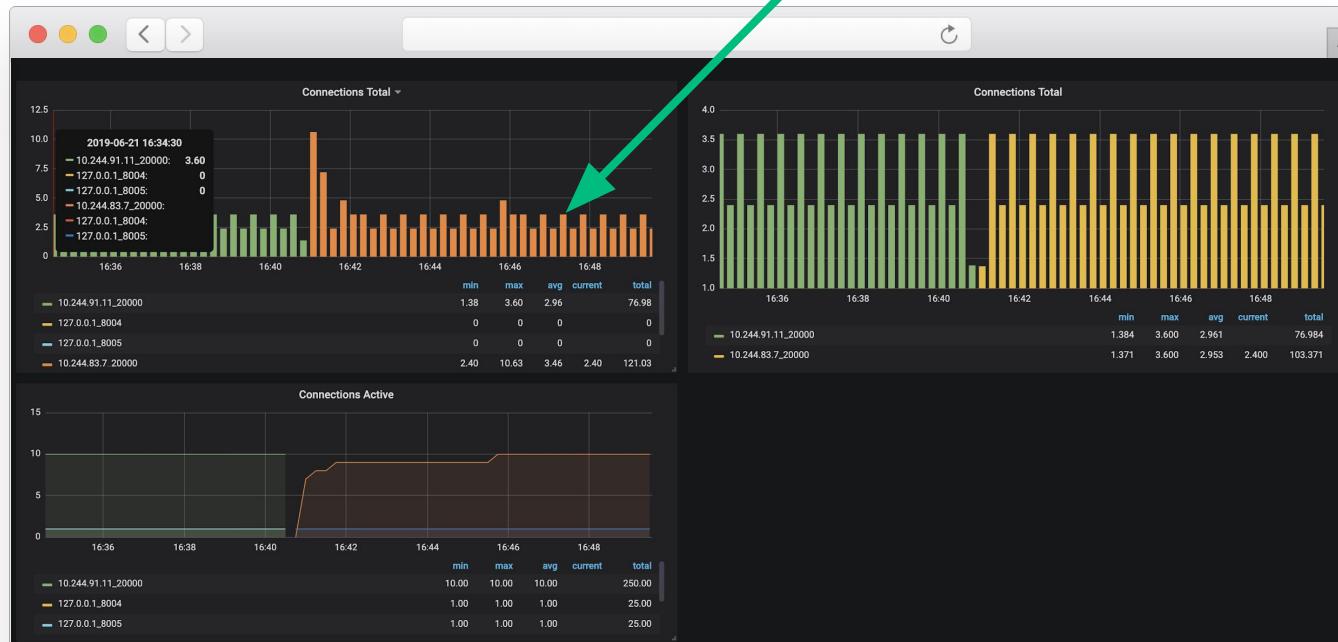
Prometheus

The screenshot shows a dark-themed code editor window with three horizontal dots at the top left and a "CODE EDITOR" label at the top right. The editor contains the following three Prometheus queries:

```
# The number of established connections to emojify-api-v2 over 30 seconds.  
increase(emoji_listener_downstream_cx_total{local_cluster="emojify-api-v2"}[30s])  
  
# The number of destroyed connections to emojify-api-v2 over 30 seconds.  
increase(emoji_listener_downstream_cx_destroy{local_cluster="emojify-api-v2"}[30s])  
  
# The current number of active connections to emojify-api-v2.  
emoji_listener_downstream_cx_active{local_cluster="emojify-api-v2"}
```

Total connections

Grafana





Key diagnostics metrics

Listener

Every listener has a statistics tree rooted at <prefix>.listener.<address>. with the following statistics:

ssl.fail_verify_no_cert	Counter	Total TLS connections that failed because of missing client certificate
ssl.connection_error	Counter	Total TLS connection errors not including failed certificate verifications
ssl.fail_verify_error	Counter	Total TLS connections that failed CA verification
ssl.fail_verify_san	Counter	Total TLS connections that failed SAN verification
downstream_pre_cx_timeout	Counter	Sockets that timed out during listener filter processing
downstream_pre_cx_active	Gauge	Sockets currently undergoing listener filter processing
downstream_cx_length_ms	Histogram	Connection length milliseconds



—

Requests HTTP / GRPC

Metrics



Key metrics

Listener - Requests HTTP

Every listener has a statistics tree rooted at <prefix>.http.<address>. with the following statistics:

downstream_rq_1xx	Counter	Total 1xx responses
downstream_rq_2xx	Counter	Total 2xx responses
downstream_rq_3xx	Counter	Total 3xx responses
downstream_rq_4xx	Counter	Total 4xx responses
downstream_rq_5xx	Counter	Total 5xx responses
downstream_rq_ws_on_non_ws_route	Counter	Total WebSocket upgrade requests rejected by non WebSocket routes
downstream_rq_time	Histogram	Total time for request and response (milliseconds)
downstream_rq_timeout	Counter	Total requests closed due to a timeout on the request path



Key metrics

Listener - Requests HTTP

Every listener has a statistics tree rooted at <prefix>.http.<address>. with the following statistics:

downstream_rq_total	Counter	Total requests
downstream_rq_http1_total	Counter	Total HTTP/1.1 requests
downstream_rq_http2_total	Counter	Total HTTP/2 requests
downstream_rq_too_large	Counter	Total requests resulting in a 413 due to buffering an overly large body
downstream_rq_completed	Counter	Total requests that resulted in a response (e.g. does not include aborted requests)



Metrics queries

Prometheus

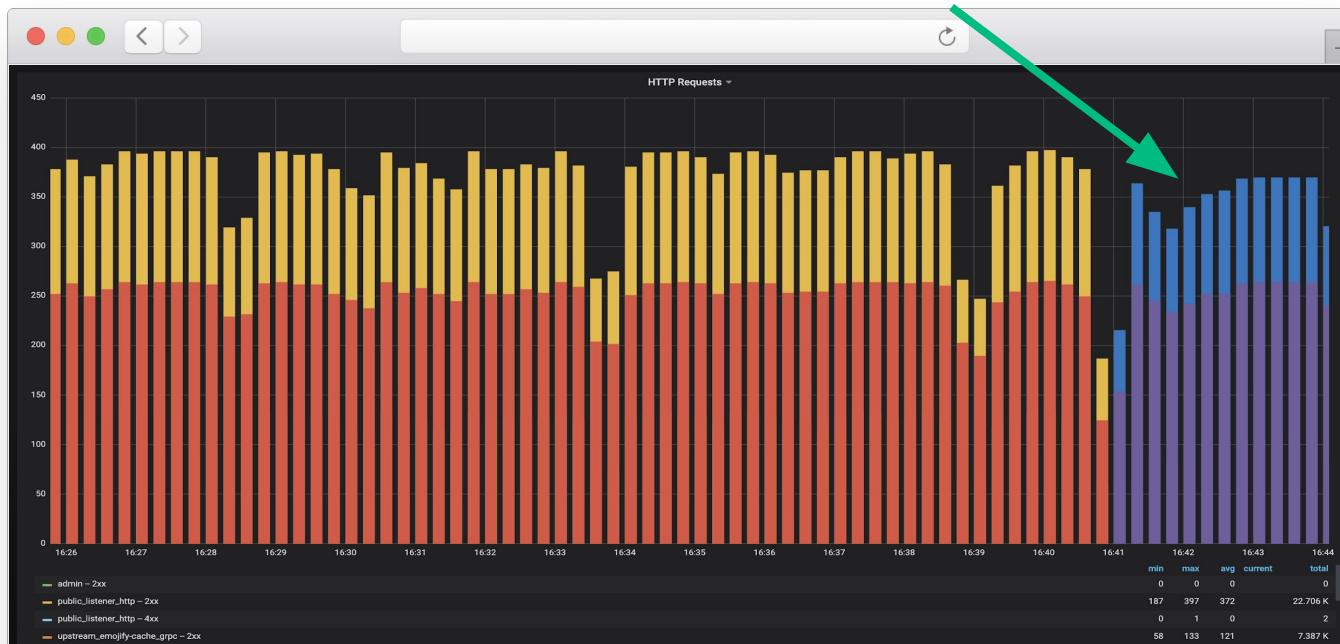
```
# The number of requests to emojify-api-v2 over 30 seconds which did not result in an error
increase(envoy_http_downstream_rq_xx{
    local_cluster="emojify-api-v2",
    envoy_response_code_class!="5"
}[30s]

# The number of requests to emojify-api-v2 over 30 seconds which resulted in an error
increase(envoy_http_downstream_rq_xx{
    local_cluster="emojify-api-v2",
    envoy_response_code_class="5"
}[30s])
```

Total Requests - all listeners for a proxy

Grafana

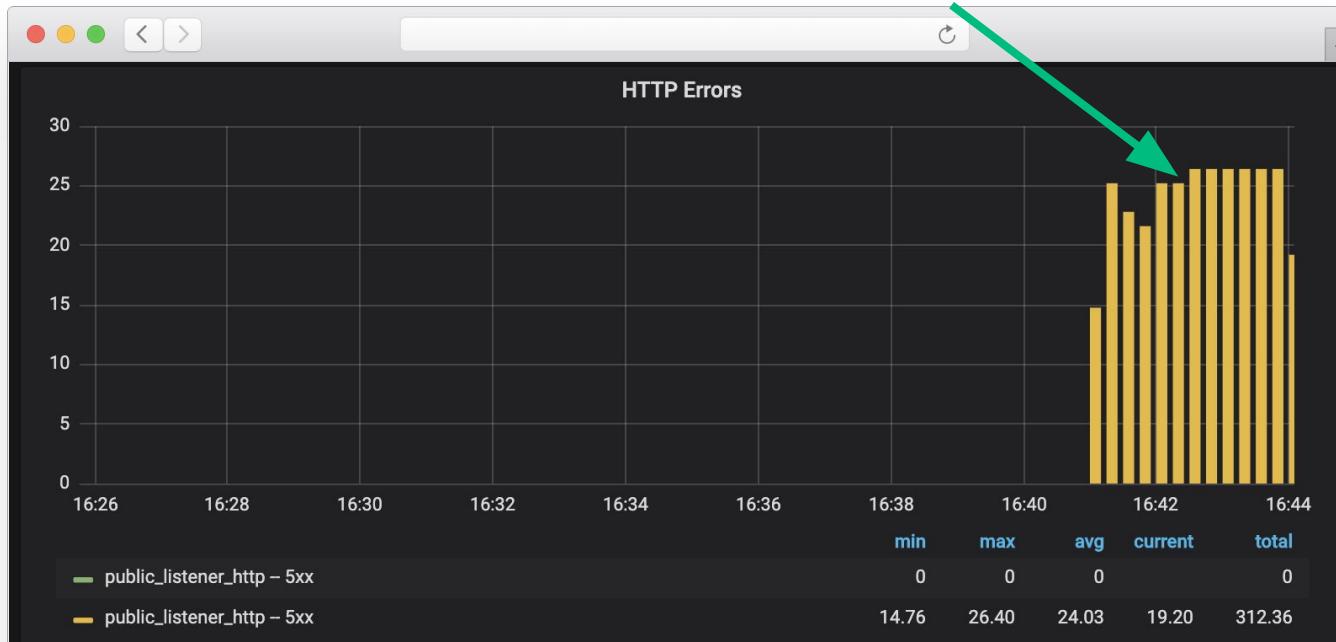
New pod started



Request Errors

Grafana

New pod started





Metrics queries - Timing

Prometheus

The screenshot shows a dark-themed Prometheus code editor window. At the top right, there's a "CODE EDITOR" button and three small circular icons. The main area contains the following Prometheus query:

```
# Upstream Timing
sum(envoy_cluster_upstream_rq_time{
    envoy_cluster_name=~"cluster_emojiify_api_v2_sidecar_proxy.*"
} > 0) by (quantile)

sum(rate(envoy_cluster_external_upstream_rq_time_sum{
    envoy_cluster_name=~"cluster_emojiify_api_v2_sidecar_proxy.*"
}[30s])) / sum(rate(envoy_cluster_external_upstream_rq_time_count{
    envoy_cluster_name=~"cluster_emojiify_api_v2_sidecar_proxy.*"
}[30s]))
```

Request Time

Grafana





Key metrics

Listener - Requests gRPC

The filter emits statistics in the cluster.<route target cluster>.grpc.namespaces

<grpc service>.<grpc method>.success	Counter	Total successful service/method calls
<grpc service>.<grpc method>.failure	Counter	Total failed service/method calls
<grpc service>.<grpc method>.total	Counter	Total service/method calls

- GRPC does **not** use HTTP status codes
- Status Codes are part of the **Protocol** and are reported as **individual** metrics

<https://www.envoyproxy.io/docs/envoy/latest/configuration/listeners/stats#listener-manager>



gRPC Bridge Filter

Configuration

The screenshot shows a dark-themed code editor window titled "CODE EDITOR". In the top left corner, there are three small circular icons. On the far left, there is a green button with a white eye icon. The main area contains JSON-style configuration code:

```
"filter_chains": [
  {
    "filters": [
      {
        "name": "envoy.http_connection_manager",
        "config": {
          "http_filters": [
            {
              "name": "envoy.grpc_http1_bridge",
              "config": {}
            },
            {
              "name": "envoy.router"
            }
          ]
        }
      }
    ]
}
```

A green rectangular box highlights the entire configuration for the "envoy.grpc_http1_bridge" filter, starting from its name and extending down to its "config" block.



Metrics queries

Prometheus

The screenshot shows a dark-themed Prometheus code editor interface. At the top, there are three small circular icons on the left and the text "CODE EDITOR" on the right. The main area contains three distinct Prometheus query blocks:

```
# GRPC no errors - Status Code 0
sum(increase(envoy_cluster_grpc_0{
    label_app="emojify-cache"
} [30s])) by (envoy_grpc_bridge_method)

# GRPC no errors - Status Code 5
sum(increase(envoy_cluster_grpc_5{
    label_app="emojify-cache"
} [30s])) by (envoy_grpc_bridge_method)

# gRPC Errors
sum(increase(envoy_cluster_grpc_failure{
    label_app="emojify-cache"
} [30s])) by (envoy_grpc_bridge_method)
```

gRPC - Success

Grafana

Methods:

Put

Get

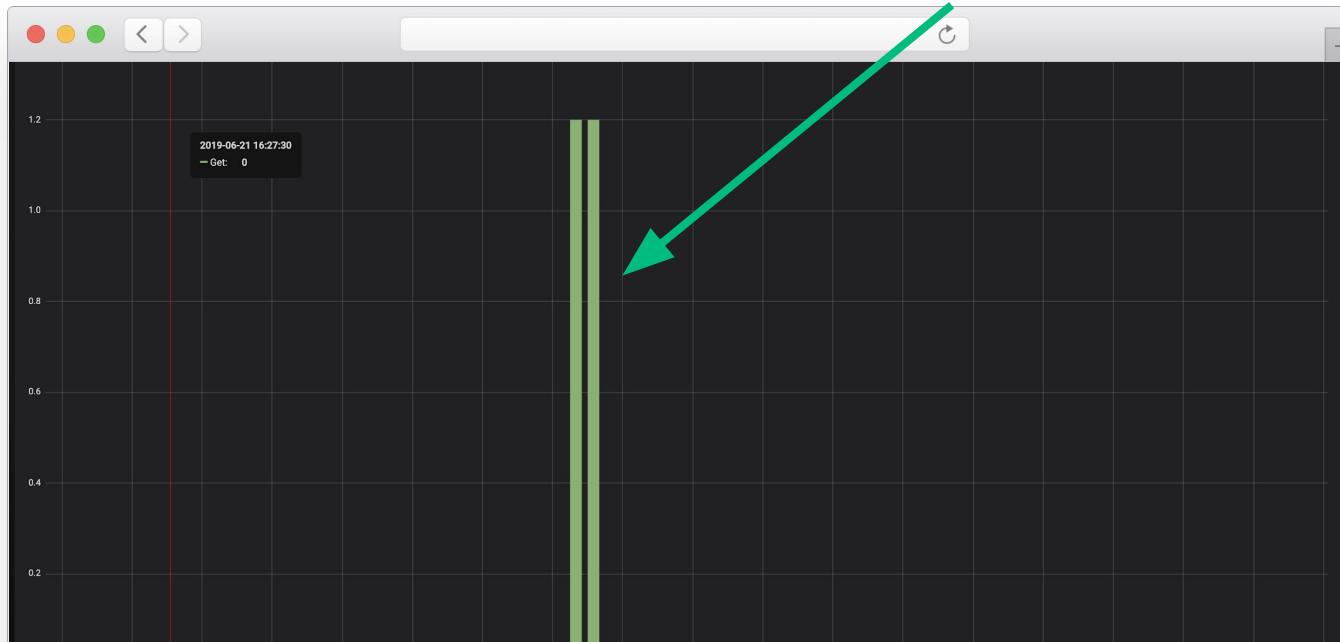
Exists



gRPC Error

Grafana

HTTP Response 5xx





—

Clusters

Metrics



Key metrics

Cluster

Every listener has a statistics tree rooted at `<prefix>.http.<address>`. with the following statistics:

upstream_rq_timeout	Counter	Total requests that timed out waiting for a response
upstream_rq_per_try_timeout	Counter	Total requests that hit the per try timeout
upstream_rq_retry	Counter	Total request retries
upstream_rq_retry_success	Counter	Total request retry successes
ejections_active	Counter	Number of currently ejected hosts



Metrics queries

Prometheus



The screenshot shows a dark-themed code editor window for Prometheus metrics queries. At the top, there are three circular tabs. On the right side of the header, the text "CODE EDITOR" is visible. The code itself is written in Prometheus query language (QL) and includes three sections: "# Retries", "# Timeouts", and "# Outlier Ejection".

```
# Retries
sum(increase(envoy_cluster_upstream_rq_retry{envoy_cluster_name=~"cluster_emojify_api_v2_sidecar_proxy.*"}[30s]))

# Timeouts
sum(increase(envoy_cluster_upstream_rq_timeout{envoy_cluster_name=~"cluster_emojify_api_v2_sidecar_proxy.*"}[30s]))
sum(increase(envoy_cluster_upstream_rq_per_try_timeout{envoy_cluster_name=~"cluster_emojify_api_v2_sidecar_proxy.*"}[30s]))

# Outlier Ejection
sum(envoy_cluster_outlier_detection_ejections_active{envoy_cluster_name=~"cluster_emojify_api_v2_sidecar_proxy.*"})
```

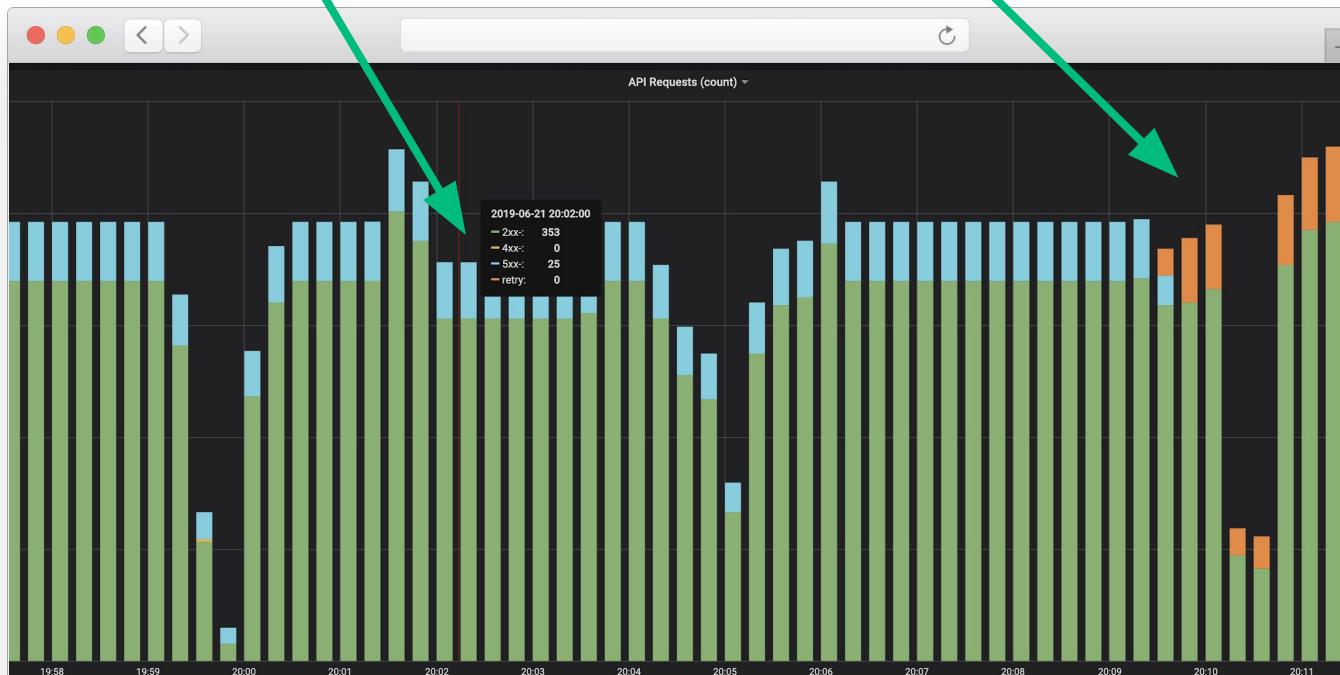
Retries

Grafana



Service Errors

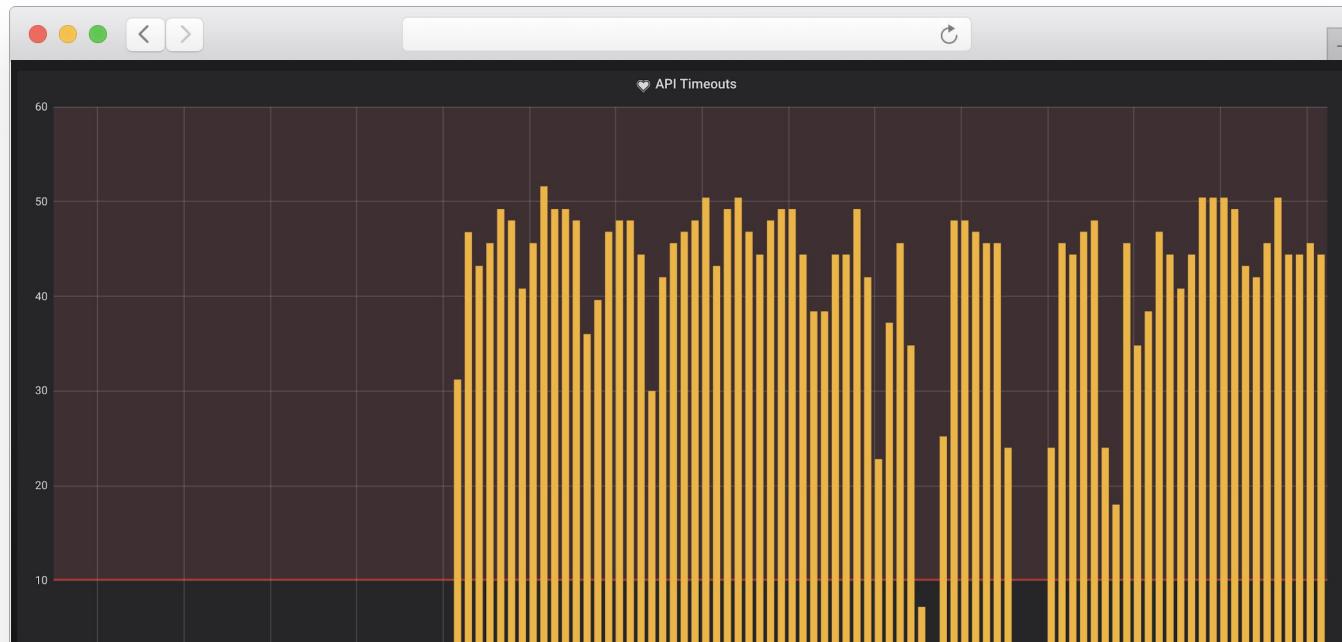
Retry Applied: No errors to user





Timeouts

Grafana



Outlier Ejection

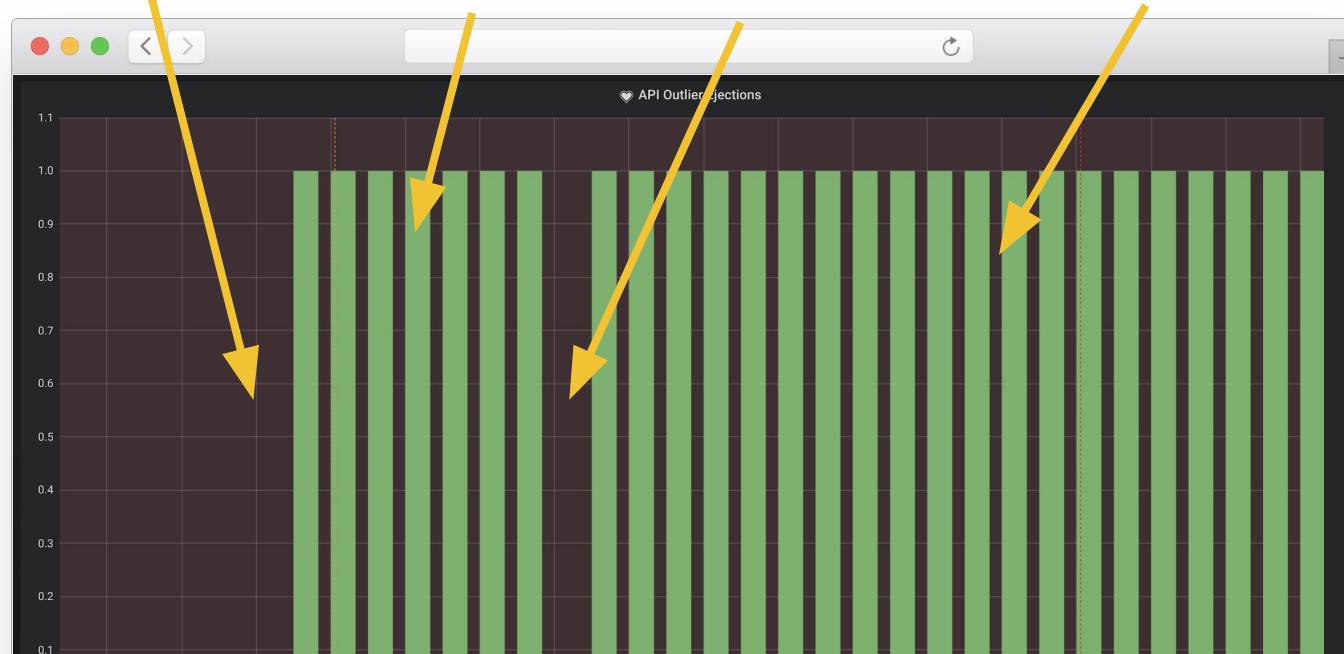
Grafana

New pod started
Constant errors

After a fixed number of
consecutive errors endpoint
removed from cluster

Envoy retries failing
endpoint

Ejection interval
increases





—

AuthZ

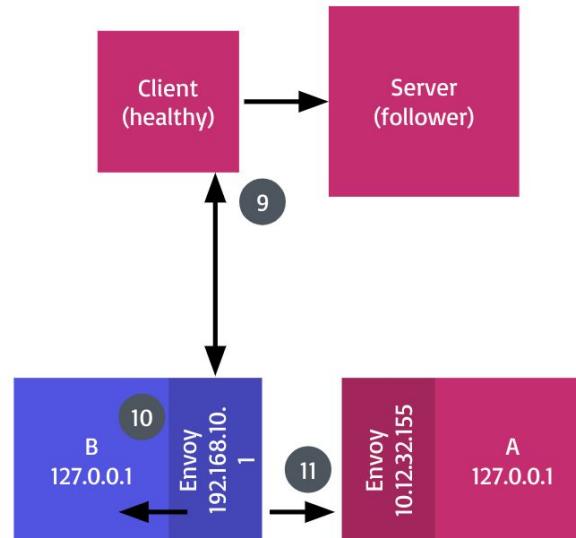
Metrics

Control plane

AuthZ

External Authorization

1. Envoy validates that the connections is allowed by calling the ext_authz filters api (once per new connection).
2. If allowed the request is passed to the upstream service.
3. Send the response to the caller.





Control plane

AuthZ

- External authorization API is normally called when establishing a new connection to an upstream.
- Failed authorization is an indication of a failing control plane, misconfiguration of security policy, or malicious activity.



Key AuthZ metrics

AuthZ

The network filter outputs statistics in the config.ext_authz namespace, with the following statistics:

total	Counter	Total responses from the filter.
error	Counter	Total errors contacting the external service.
denied	Counter	Total responses from the authorizations service that were to deny the traffic.
failure_mode_allowed	Counter	Total requests that were error(s) but were allowed through because of failure_mode_allow set to true.
ok	Counter	Total responses from the authorization service that were to allow the traffic.
cx_closed	Counter	Total connections that were closed.
active	Gauge	Total currently active requests in transit to the authorization service.

https://www.envoyproxy.io/docs/envoy/latest/configuration/network_filters/ext_authz_filter.html



Metrics queries

Prometheus

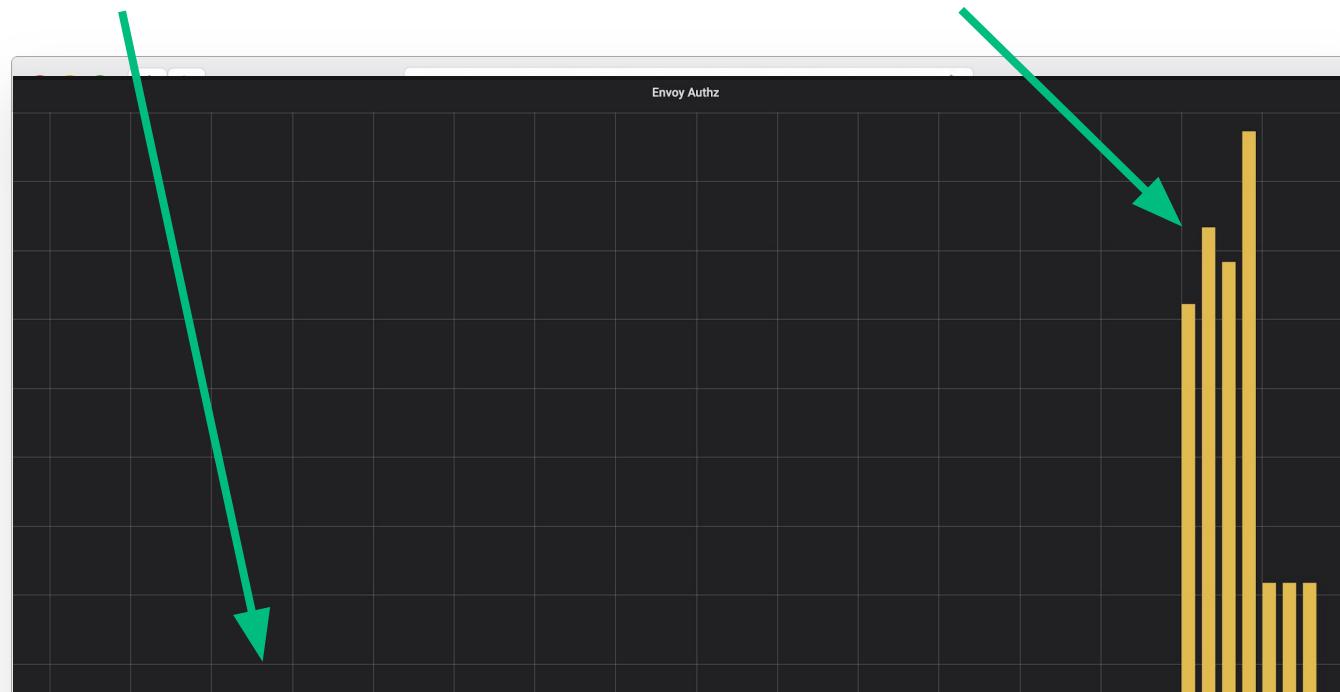
The image shows a dark-themed code editor window with three circular navigation dots at the top left. At the top right, the text "CODE EDITOR" is displayed. The editor contains two lines of Prometheus query text:

```
# Successful AuthZ  
increase(envoy_ext_authz_connect_authz_ok{local_cluster="emojify-api-v2"}[1m])  
  
# AuthZ Denied  
increase(envoy_ext_authz_connect_authz_denied{local_cluster="emojify-api-v2"}[1m])
```

AuthZ OK

Grafana

Cached Authorization - no metrics



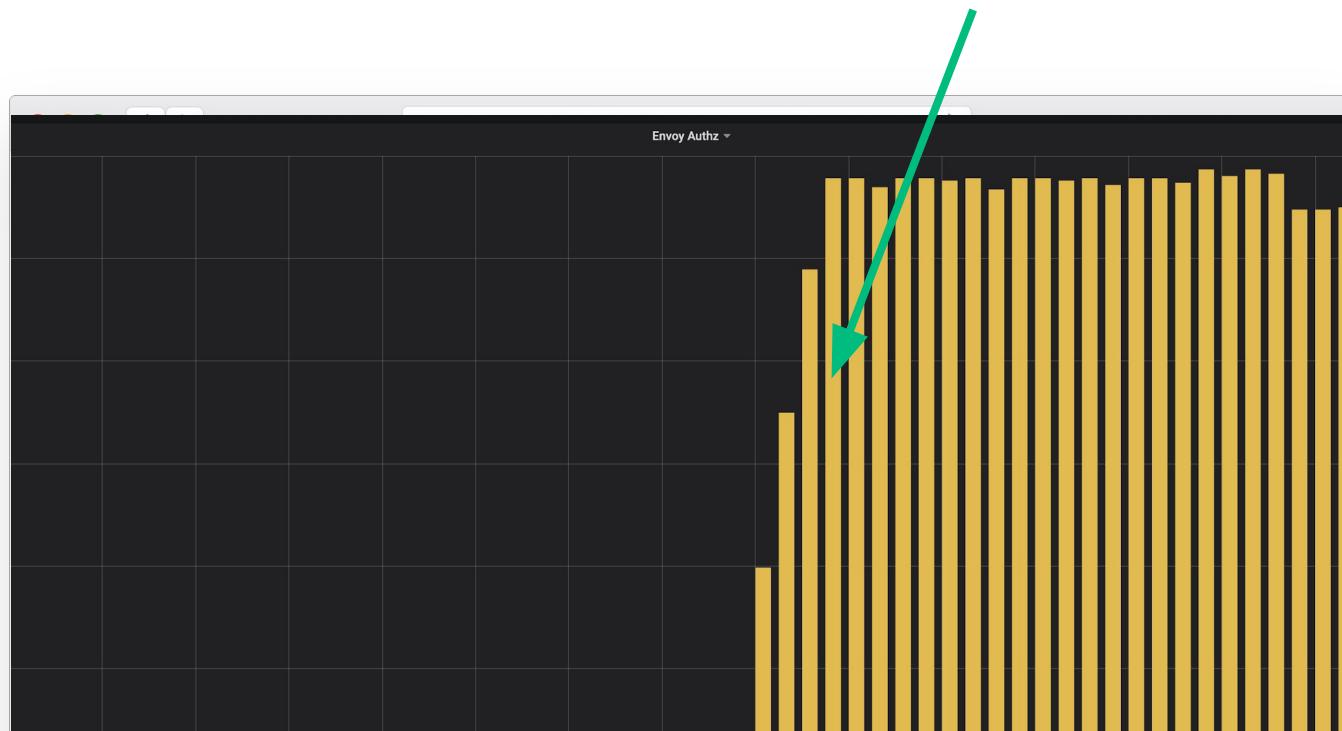
New pod started



AuthZ Failed

Grafana

AuthZ failure, either:
Misconfiguration or Attack





—

Tracing



What is tracing?

Tracing

Distributed tracing, also called distributed request tracing, is a **method used to profile and monitor applications**, especially those built using a microservices architecture. **Distributed tracing helps pinpoint where failures occur and what causes poor performance.**

<https://opentracing.io/docs/overview/what-is-tracing/>



—

Configuration

Tracing



Tracing Cluster

Configuration

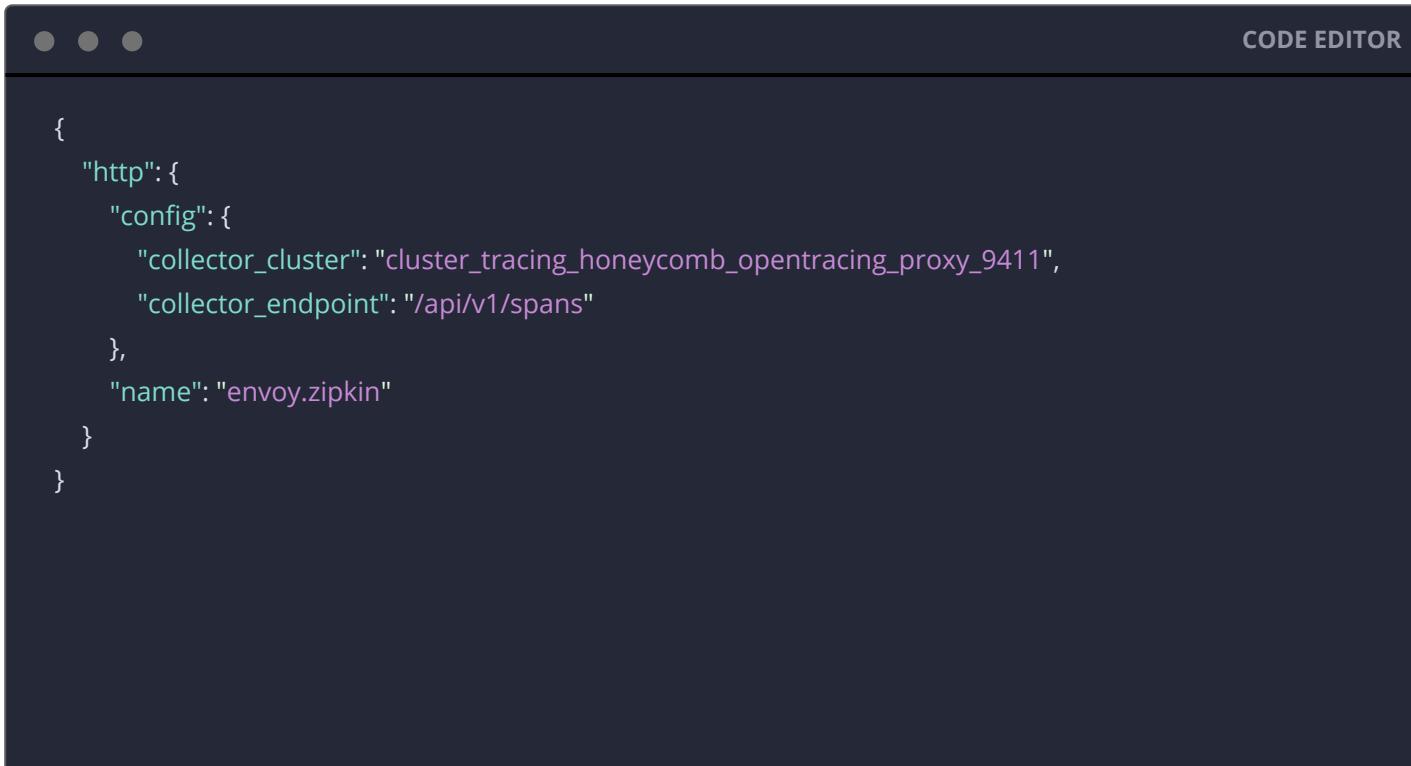
A screenshot of a dark-themed code editor window titled "CODE EDITOR". The window shows a single JSON configuration file. The code is color-coded, with strings in purple and various keys and values in different shades of green and blue. The JSON structure defines a "load_assignment" object with a "cluster_name", "endpoints", and "lb_endpoints" array. Each endpoint has an "address" object with "socket_address", "port_value", and "protocol" fields. The "name" field at the bottom is also purple.

```
"load_assignment": {  
    "cluster_name": "cluster_tracing_honeycomb_opentracing_proxy_9411",  
    "endpoints": [ {  
        "lb_endpoints": [ {  
            "endpoint": {  
                "address": {  
                    "socket_address": {  
                        "address": "honeycomb-opentracing-proxy",  
                        "port_value": 9411,  
                        "protocol": "TCP"  
                    } } }  
            } ]  
        } ],  
    "name": "cluster_tracing_honeycomb_opentracing_proxy_9411"  
}
```



Tracing Configuration

Configuration



A dark-themed code editor interface showing a JSON configuration file. The editor has three circular navigation dots at the top left and a "CODE EDITOR" label at the top right. The code itself is a single-line JSON object with nested properties for "http" and "name".

```
{  
  "http": {  
    "config": {  
      "collector_cluster": "cluster_tracing_honeycomb_opentracing_proxy_9411",  
      "collector_endpoint": "/api/v1/spans"  
    },  
    "name": "envoy.zipkin"  
  }  
}
```

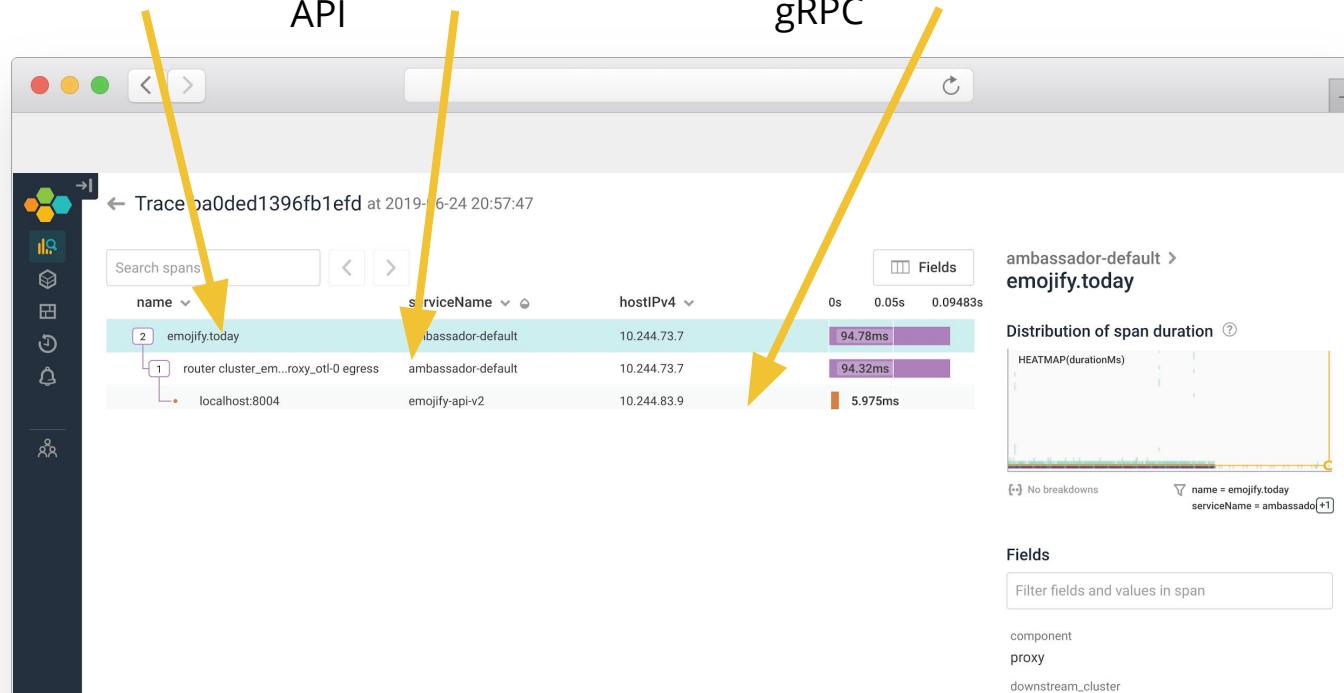
Trace - HTTP Post

honeycombi.io

Public Ingress

Route upstream to
API

External upstream
gRPC





Handling tracing spans

Tracing



Adding Headers

HTTP

The image shows a screenshot of a code editor window with a dark theme. The title bar says "CODE EDITOR". The code in the editor is written in Go and defines a function for adding specific headers to an http.Header.

```
var otHeaders = []string{
    "x-request-id",
    "x-b3-traceid",
    "x-b3-spanid",
    "x-b3-parentspanid",
    "x-b3-sampled",
    "x-b3-flags",
    "x-ot-span-context"}
var headers http.Header
for _, h := range otHeaders {
    if v := r.Header.Get(h); len(v) > 0 { headers.Add(h, v) }
}

return headers
```



Adding Headers

HTTP

```
headers := createHeadersFromRequest(r)

req, _ := http.NewRequest("GET", "http://localhost:8004", nil)
req.Header = headers

resp, err := http.DefaultClient.Do(req)
if err != nil {
    http.Error(rw, err.Error(), http.StatusInternalServerError)
    return
}
```

CODE EDITOR



Adding Headers

gRPC

The screenshot shows a dark-themed code editor window titled "CODE EDITOR". At the top left, there are three small circular icons. On the right side of the title bar, there is a "CODE EDITOR" label. The main area contains the following Go code:

```
var otHeaders = []string{
    "x-request-id",
    "x-b3-traceid",
    "x-b3-spanid",
    "x-b3-parentspanid",
    "x-b3-sampled",
    "x-b3-flags",
    "x-ot-span-context"
}

var pairs []string
for _, h := range otHeaders {
    if v := r.Header.Get(h); len(v) > 0 { pairs = append(pairs, h, v) }
}
md := metadata.Pairs(pairs...)
return metadata.NewOutgoingContext(context.Background(), md)
```



Adding Headers

gRPC

```
// create a grpc context containing the parent span metadata
ctx := createGRPCContextFromRequest(r)

resp, err := e.emojify.Create(ctx, &wrappers.StringValue{Value: u.String()})
if err != nil {
    http.Error(rw, err.Error(), http.StatusInternalServerError)
    return
}
```

CODE EDITOR



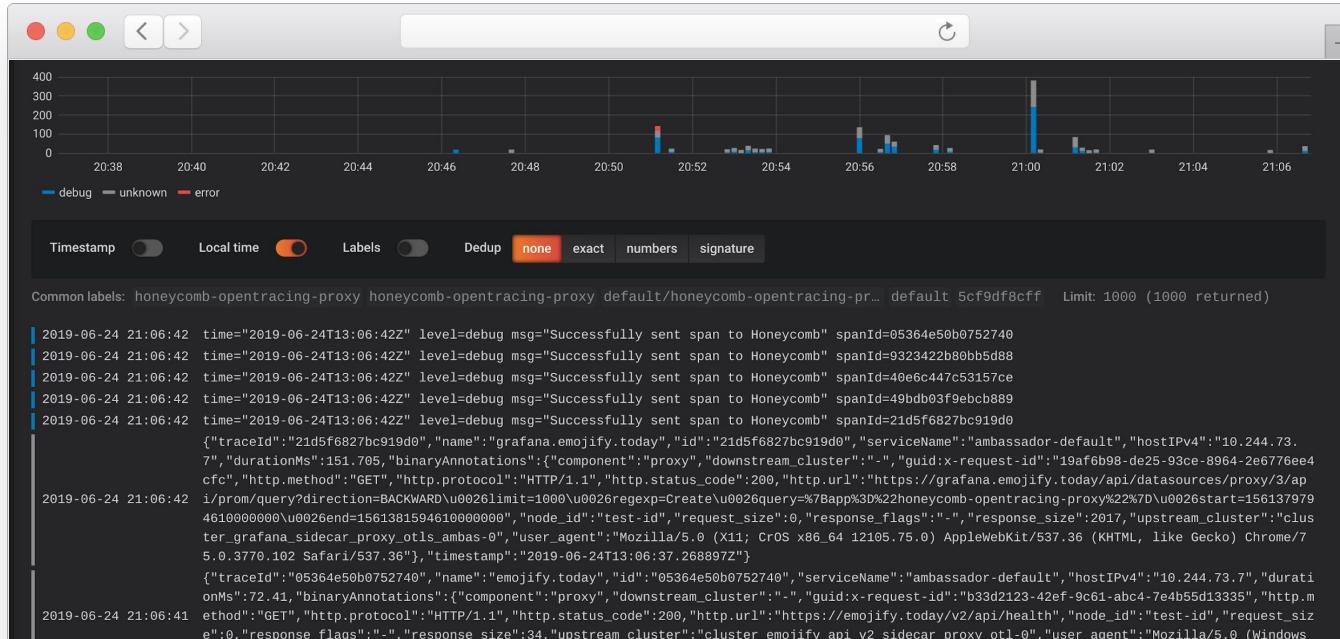
—

Logging



Logging

Loki





Thank You

nic@hashicorp.com
www.hashicorp.com