



# OPEN SOURCE SUMMIT

---

China 2019

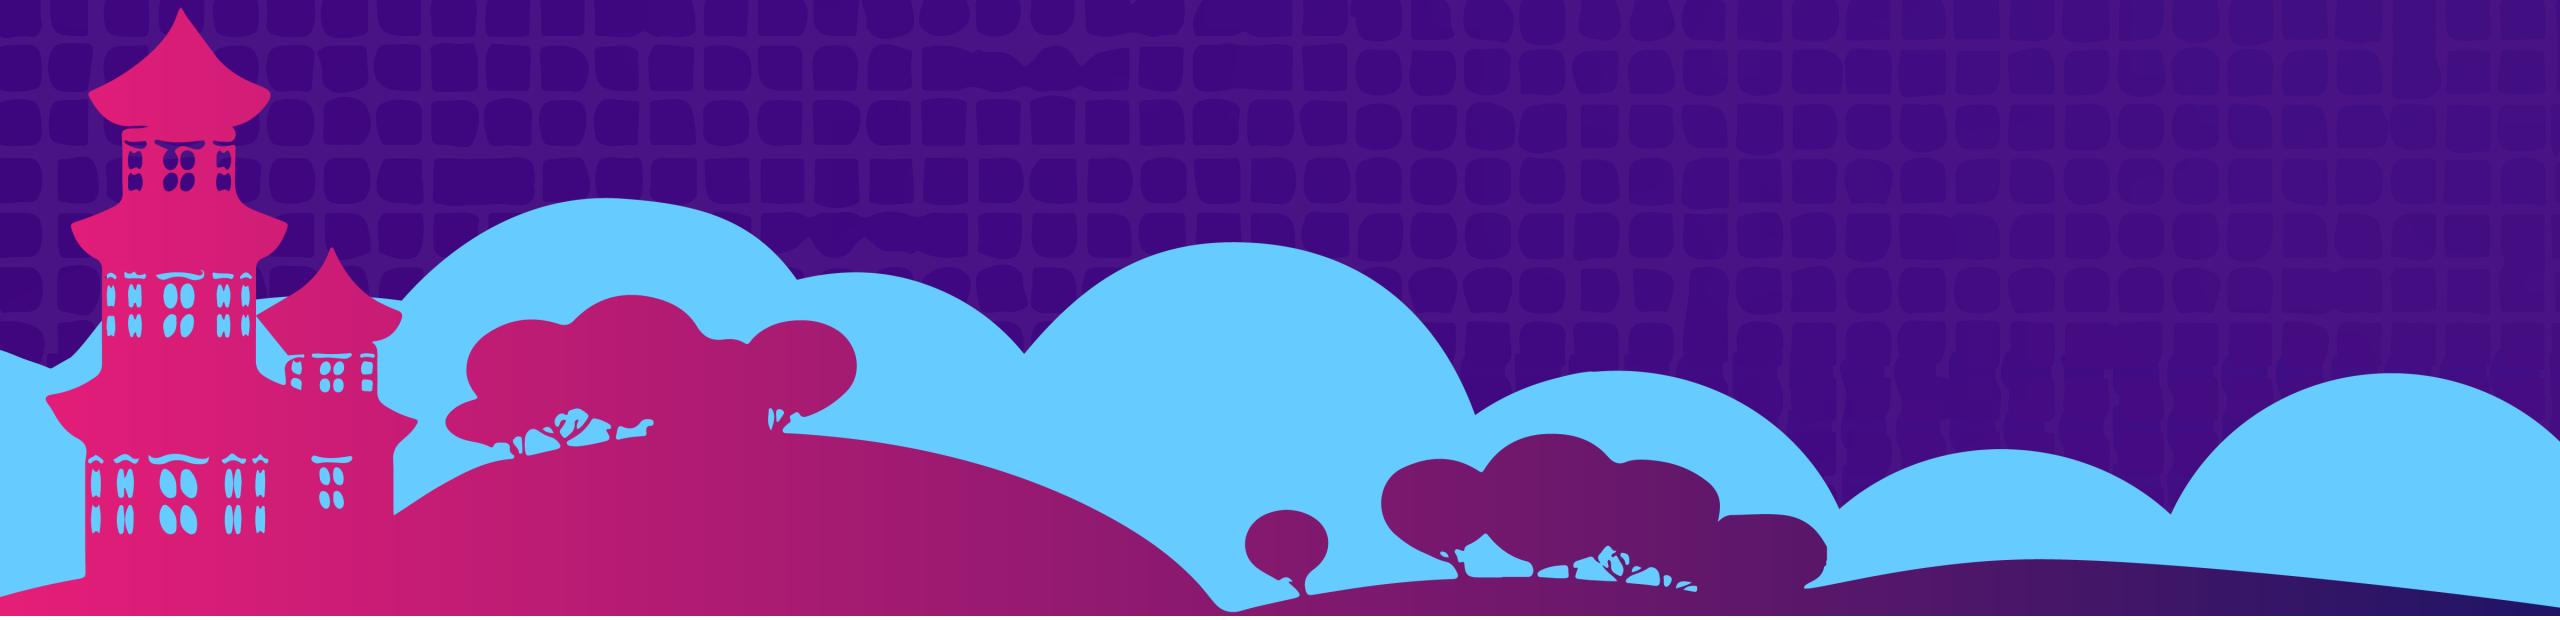
---



# Linux Kernel Live Patching

Haishuang Yan

[<yanhaishuang@cmss.chinamobile.com>](mailto:yanhaishuang@cmss.chinamobile.com)





# What is Live patching

- Live patching function for kernel
  - Apply a binary to kernel-on-line
  - Patching is done without reboot
  - No Disruption to applications
- Used for security fixes
  - Only for a small and critical issues
  - Not for major kernel updates



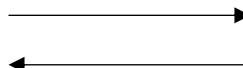
# Why use Live patching

- Many kernel security bugs to be fixed
- Kernel update = reboot
  - Disruption to applications
  - Hardware reboot failures
- High cost of downtime caused by reboot
  - Some applications can't accept 100ms downtime
- Scheduled downtime

# From Ftrace to Kpatch

- Kernel Function Tracer
- GCC profiler option: -pg -fentry
- Adds special fentry function call
  - All function call fentry
  - Fentry is a trampoline

```
<schedule>:  
callq ffffffff81a01d30 <__fentry__>  
mov %gs:0x15c40,%rax  
push %rbx  
...
```



```
<__fentry__>:  
retq  
nopl 0x0(%rax,%rax,1)  
nopw %cs:0x0(%rax,%rax,1)
```

# Dynamic Ftrace

- Can't just call the fentrys
  - Too much overhead
  - Big impact on performance(13%)
- On boot convert all the locations to NOPS

```
<schedule>:  
nopl 0x0(%rax,%rax,1) [FTRACE NOP]  
mov %gs:0x15c40,%rax  
push %rbx  
...
```

# Enable Tracing

<schedule>:

```
0xfffffffffa77fc9a0 <schedule>: nopl 0x0(%rax,%rax,1) [FTRACE NOP]
0xfffffffffa77fc9a5 <schedule+5>:    mov %gs:0x15c40,%rax
0xfffffffffa77fc9ae <schedule+14>:   push %rbx
0xfffffffffa77fc9af <schedule+15>:   mov 0x10(%rax),%rdx
0xfffffffffa77fc9b3 <schedule+19>:   test %rdx,%rdx
0xfffffffffa77fc9b6 <schedule+22>:   je  0xfffffffffa77fc9c2 <schedule+34>
0xfffffffffa77fc9b8 <schedule+24>:   cmpq $0x0,0xb90(%rax)
0xfffffffffa77fc9c0 <schedule+32>:   je  0xfffffffffa77fc9db <schedule+59>
0xfffffffffa77fc9c2 <schedule+34>:   mov %gs:0x15c40,%rbx
0xfffffffffa77fc9cb <schedule+43>:  xor %edi,%edi
0xfffffffffa77fc9cd <schedule+45>:  callq 0xfffffffffa77fc130 <__schedule>
```

# Enable Tracing

```
<schedule>:  
0xfffffffffa77fc9a0 <schedule>: callq 0xfffffffffa7a01f20 <ftrace_graph_caller>  
0xfffffffffa77fc9a5 <schedule+5>:    mov    %gs:0x15c40,%rax  
0xfffffffffa77fc9ae <schedule+14>:   push   %rbx  
0xfffffffffa77fc9af <schedule+15>:   mov    0x10(%rax),%rdx  
0xfffffffffa77fc9b3 <schedule+19>:   test   %rdx,%rdx  
0xfffffffffa77fc9b6 <schedule+22>:   je     0xfffffffffa77fc9c2 <schedule+34>  
0xfffffffffa77fc9b8 <schedule+24>:   cmpq   $0x0,0xb90(%rax)  
0xfffffffffa77fc9c0 <schedule+32>:   je     0xfffffffffa77fc9db <schedule+59>  
0xfffffffffa77fc9c2 <schedule+34>:   mov    %gs:0x15c40,%rbx  
0xfffffffffa77fc9cb <schedule+43>:  xor    %edi,%edi  
0xfffffffffa77fc9cd <schedule+45>:  callq 0xfffffffffa77fc130 <__schedule>
```

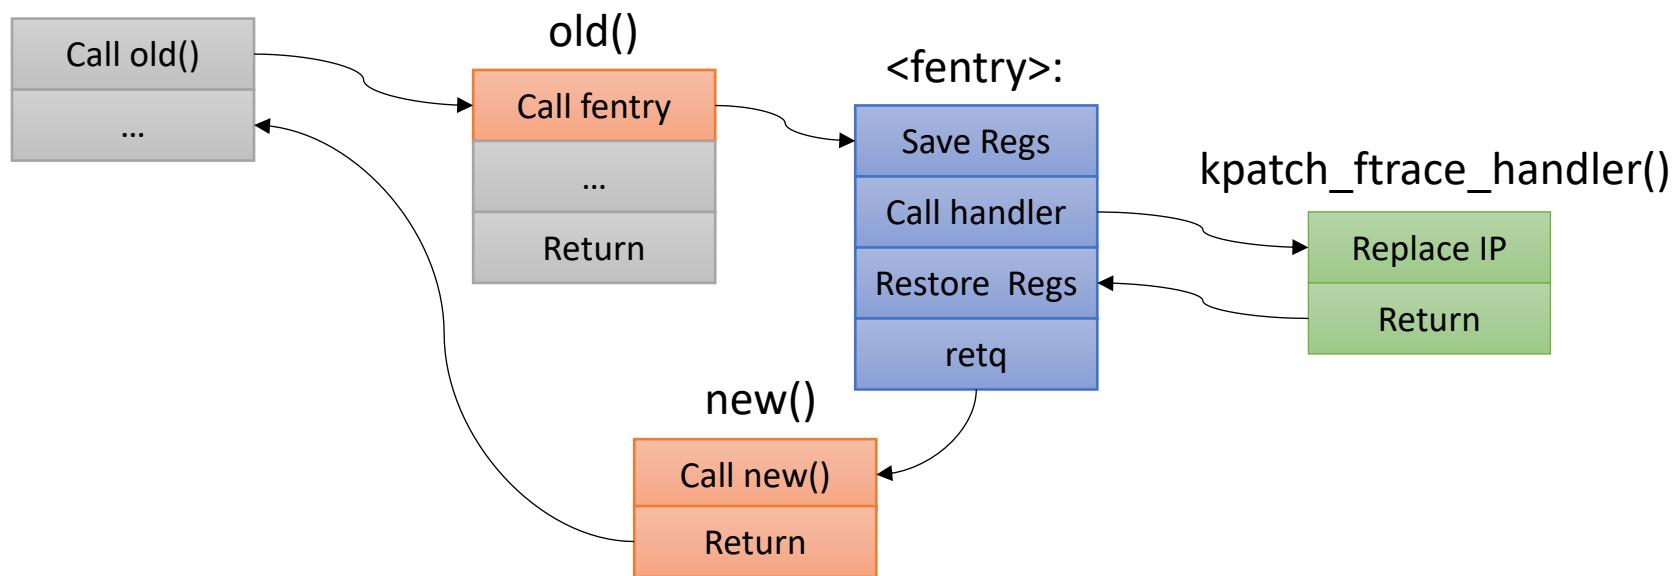
```
# echo schedule > set_ftrace_filter  
# echo function_graph > current_tracer  
# cat set_ftrace_filter  
schedule
```

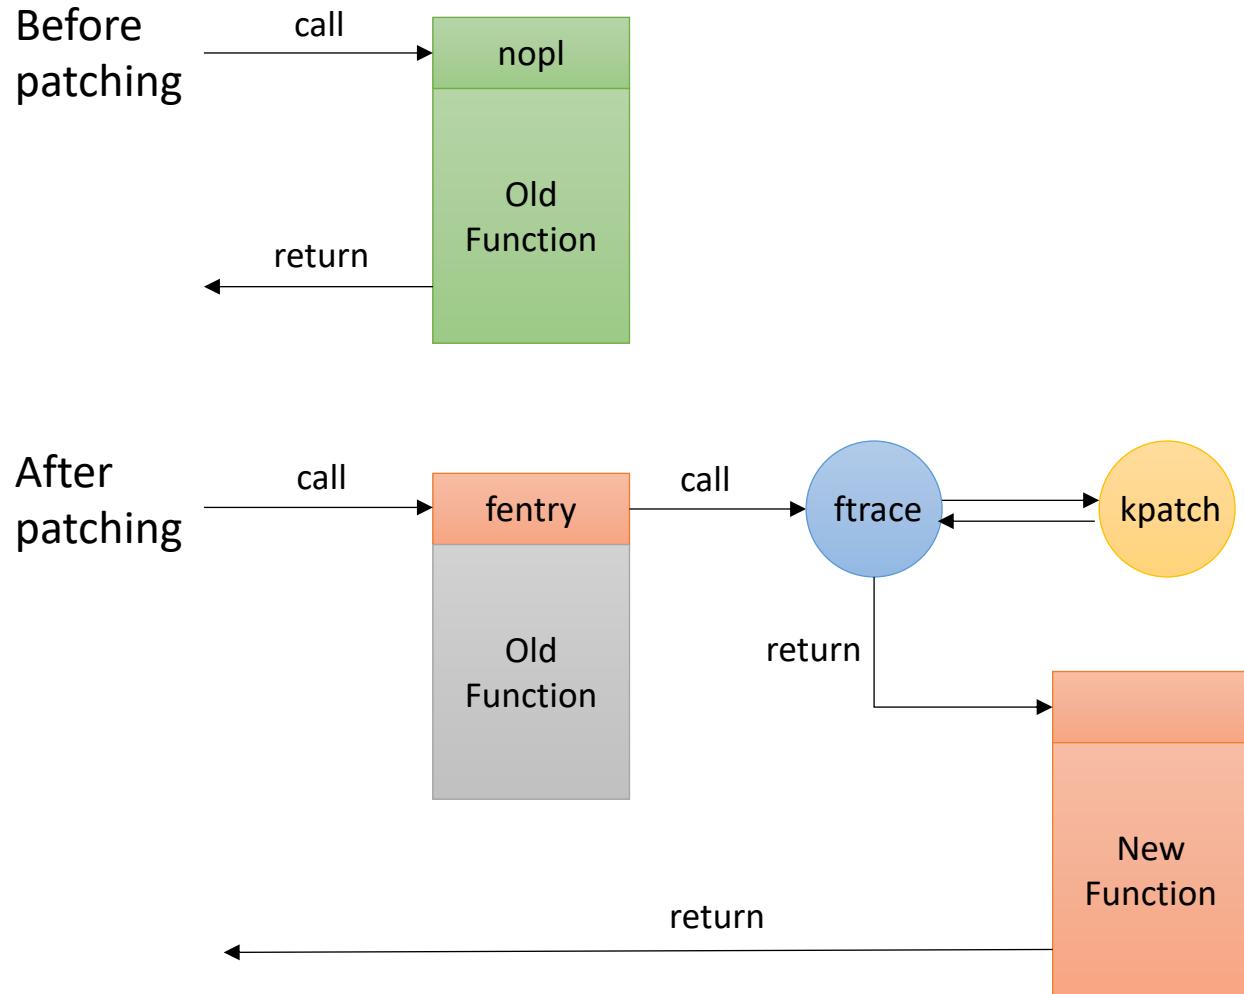
# Kernel Live Patching



- Add Ftrace Entry
- Stop\_machine()
  - Stop running processes for a while
  - Disable interrupts
- Safeness check
  - Walk through the threads and check the stack
- Enable the hook
  - Use Ftrace

- Kpatch uses Dynamic Ftrace to patch
  - Add old function to ftrace filter and register handler
  - Replace regs->ip with new function address





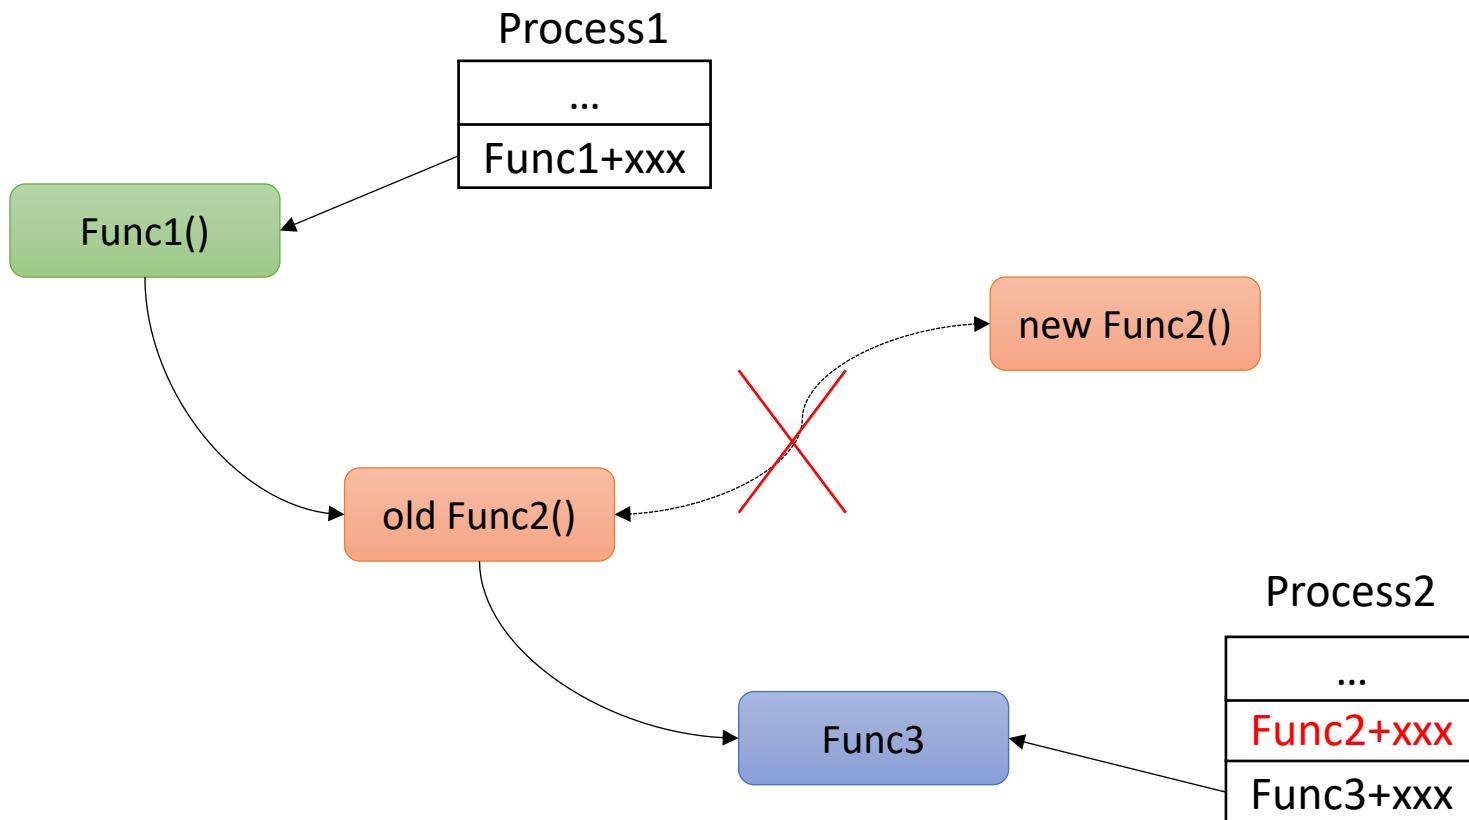


# Activeness Safety Check

- Verify activeness safety
  - Ensures none of the to-be-patched functions are on the stack of any task
  - This function is called from stop\_machine() context
  - Walk through the all Thread and check old functions on stacks
  - Verify the backtrace address on the stack

# Activeliness Safety Check

OPEN SOURCE SUMMIT  
China 2019

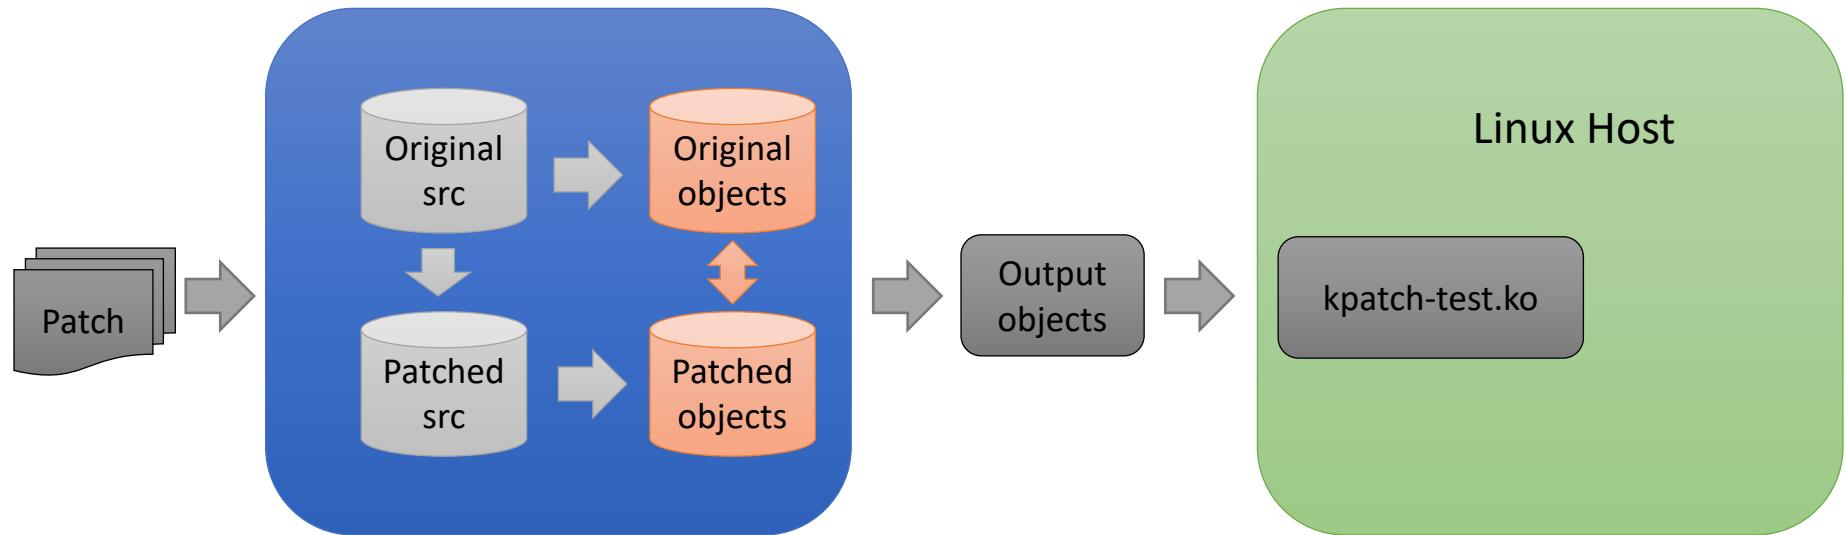


- **kpatch-build**
  - A collection of tools which convert a source diff patch to a patch module
- **Patch Module**
  - A kernel module (.ko file) which includes the replacement functions and metadata about the original functions
- **Kpatch Core Module**
  - A kernel module (.ko file) which provides an interface for the patch modules to register new functions for replacement

- kpatch-build
  - Build unstripped vmlinux for the original kernel
    - With -ffunction-sections -fdata-sections flags
  - Rebuild vmlinux for patched kernel
    - Watch for changed objects
  - Analyze each original/patched objects
    - Generate three resulting output object
  - Link all the output Objects
    - Generate the patch module

# Kpatch Tools

- Build the patch module
  - `kpatch-build test.patch -n kpatch-test`
- Patch the kernel
  - `kpatch load kpatch-test.ko`



# Demo



# Limitations

- Human safety analysis required!
- Patches which modify init functions are not supported
- Patches which modify functions that are missing a fentry call are not supported
- 80% of all CVE patches currently supported
- `stop_machine()` latency: 1ms-40ms

# Thanks

